



**MANIPAL INSTITUTE OF TECHNOLOGY**

**MANIPAL**

*(A constituent unit of MAHE, Manipal)*

**DEPARTMENT OF COMPUTER SCIENCE &  
ENGINEERING**

**CERTIFICATE**

This is to certify that Ms./Mr. ....

Reg. No.: ..... Section: ..... Roll No.: .....

has satisfactorily completed the lab exercises prescribed for OPERATING  
SYSTEMS LAB [CSE 3142] of Third Year B. Tech. Degree at MIT, Manipal, in  
the academic year 2025-2026.

Date: .....

Signature

Faculty in Charge

# CONTENTS

LAB NO.	TITLE	PAGE NO.	REMARKS
	COURSE OBJECTIVES AND OUTCOMES	v	
	EVALUATION PLAN	v	
	INSTRUCTIONS TO THE STUDENTS	vi	
1	LINUX BASIC COMMANDS, SHELL CONCEPTS AND FILE FILTERS		
2	SHELL SCRIPTING - 1		
3	SHELL SCRIPTING - 2		
4	PROCESSES AND SIGNALS		
5	FILE SYSTEM		
6	IPC - 1 PIPE, FIFO		
7	IPC -2 MESSAGE QUEUE , SHARED MEMORY		
8	IPC-3 -DEADLOCK , LOCKING SYNCHRONIZATION		
9	PROGRAMS ON THREADS		
10	MEMORY AND DATA MANAGEMENT		
11	DEADLOCK AND DISK MANAGEMENT		
12	WORKING WITH REGULAR FILES AND DIRECTORY STRUCTURES		
	REFERENCES		







## **Course Objectives**

- Illustrate and explore system calls related to Linux operating system.
- Learn process management and thread programming concepts which include scheduling algorithms and inter process communication.
- Understand the working of memory management schemes, disk scheduling algorithms, and page replacement algorithms through simulation.

## **Course Outcomes**

At the end of this course, students will be able to:

- Execute Linux commands, shell scripting using appropriate Linux system calls.
- Design thread programming, simulate process management and inter process communication techniques.
- Implement the memory management, disk scheduling and page replacement algorithms.

## **Evaluation Plan**

- Internal Assessment Marks : 60%
  - ✓ Continuous Observation component (for each experiment):  $2 \times 12 = 24$  marks
  - ✓ Continuous Execution component (for each experiment) :  $(1+2) \times 12 = 36$  marks
  - ✓ Total marks of the 12 experiments will sum up to 60
- End semester assessment of 2-hours duration: 40 Marks

## **INSTRUCTIONS TO THE STUDENTS**

### **Pre-Lab Session Instructions**

1. Students should carry the Class notes, Lab Manual and the required stationery to every lab session.
2. Be in time and follow the Instructions from Lab Instructors.
3. Must Sign in the log register provided.
4. Make sure to occupy the allotted seat and answer the attendance.
5. Adhere to the rules and maintain the decorum.

### **In-Lab Session Instructions**

- Follow the instructions on the allotted exercises given in Lab Manual.
- Show the program and results to the instructors on completion of experiments.
- On receiving approval from the instructor, copy the program and results in the Lab record.
- Prescribed textbooks and class notes can be kept ready for reference if required.

### **General Instructions for the exercises in Lab**

- The programs should meet the following criteria:
  - Programs should be interactive with appropriate prompt messages, error messages if any, and descriptive messages for outputs.
  - Programs are properly indented and comments should be given whenever it is required.
  - Use meaningful names for variables and procedures.
- Plagiarism (copying from others) is strictly prohibited and would invite severe penalty during evaluation.
- The exercises for each week are divided under three sets:
  - Solved exercise
  - Lab exercises - to be completed during lab hours
  - Additional Exercises - to be completed outside the lab or in the lab to enhance the skill.

- In case a student misses a lab class, he/she must ensure that the experiment is completed at students end or in a repetition class (if available) with the permission of the faculty concerned but credit will be given only to one day's experiment(s).
- Questions for lab tests and examination are not necessarily limited to the questions in the manual, but may involve some variations and/or combinations of the questions.
- A sample note preparation is given later in the manual as a model for observation.

### **THE STUDENTS SHOULD NOT**

- Carry mobile phones while working with computer.
- Go out of the lab without permission.



## **LINUX BASIC COMMANDS, SHELL CONCEPTS AND FILE FILTERS**

### **Objectives:**

In this lab, the student will be able to:

- Learn to Linux basic commands
- Understand the working of commands and important shell concepts and file filters
- Write and execute basic commands in a Shell

**shell** a utility program that enables the user to interact with the Linux operating system. Commands entered by the user are passed by the shell to the operating system for execution. The results are then passed back by the shell and displayed on the user's display. There are several shells available like Bourne shell, C shell, Korn shell, etc. Each shell differs from the other in Command interpretation. The most popular shell is bash.

**shell prompt** a character at the start of the command line which indicates that the shell is ready to receive the commands. The character is usually a '%' (percentage sign) or a '\$' (dollar sign).

For. e.g.

Last login : Thu April 11 06:45:23

\$ \_ (This is the shell prompt, the cursor shown by the \_ character).

Linux commands are executable binary files located in directories with the name bin (for binary). Many of the commands that are generally used are located in the directory /usr/bin.

**echo** is a command for displaying any string in the command prompt.

For e.g. \$ echo "Welcome to MIT Manipal"

Environment variables: Shell has built in variables which are called environment variables. For e.g. the user who has logged in can be known by typing

**\$echo \$USER**

The above will display the current user's name.

When the command name is entered, the shell checks for the location of the command in each directory in the PATH environment variable. If the command is found in any of the directories mentioned in PATH, then it will execute. If not found, will give a message "Command not found".

## COMMONLY USED LINUX COMMANDS

**who:** Unix is a system that can be concurrently used by multiple users and to know the users who are using the system can be known by a **who** command. For e.g. Current users are kumar, vipul and raghav. These are the user ids of the current users.

```
$ who [Enter] kumar pts/10 May  
1 09.32 vipul pts/4 May 1  
09.32 raghav pts/5 May 1 09.32
```

The first columns indicates the user name of the user, second column indicates the terminal name and the third column indicates the login time. To know the user who has invoked the command can be known by the following command. For e.g. if kumar is the user who has typed the who command above then, \$ **who am i** [Enter] kumar pts/10 May 1 09.32 **ls:** UNIX system has a large number of files that control its functioning and users also create files on their own. These files are stored in separate folders called directories. We can list the names of the files available in this directory with **ls** command. The list is displayed in the order of creation of files.

```
$ ls [Enter]  
README  
chap01  
chap02  
chap03  
helpdir  
progs
```

In the above output, **ls** displays a list of six files. We can also list specific files or directories by specifying the file name or directory names. In this we can use regular expressions.

For e.g. to list all files beginning with chap we can use the following command.

```
$ ls chap* [Enter] chap01  
chap02  
chap03
```

To list further detailed information we can use **ls -l** command, where **-l** is an option between the command and filenames. The details include, file type, file or directory access permissions, number of links, owner name, group name, file or directory size, modification time and name of file or directory.

```
$ ls -l chap* [Enter]
```

-rw-r- - r- -l kumar users	5670 Apr 3 09.30 chap01
-rw-r- - r- -l kumar users	5670 Feb 23 09.30 chap02
-rw-r- - r- -l kumar users	5670 Apr 30 09.34 chap03

The argument beginning with hyphen is known as option. The main feature of option is it starts with hyphen. The command **ls** prints the columnar list of files and directories. With the **-l** option it displays all the information as shown above.

General syntax of **ls** command:

***ls -[options][file list][directory list]***

In Linux, file names beginning with period are hidden files, are not normally displayed in **ls** command. To display all files, including the hidden ones, use option **-a** in **ls** command as shown below:

**\$ ls -a**

**\$ ls /** will display the name of the files and sub-directories under the root directory.

**pwd:** This command gives the present working directory where the user is currently located.

**\$ pwd**

/home/kumar/pis

**cd:** To move around in the file system use **cd** (change directory) command. When used with argument, it changes the current directory to the directory specified as argument, for instance:

**\$ pwd**

/home/kumar

**\$ cd progs**

**\$ pwd**

\$ /home/kumar/progs

**cd ..** : To change the working directory to the parent of the current directory we need to use **\$ cd ..**

.. (double dot) indicates parent directory. A single dot indicates current directory.

**cat:** **cat** is a multipurpose command. Using this we can display a file, create a file as well as concatenate files vertically.

```
$ cat > filename[Enter] cat >  
os.txt
```

Welcome to Manipal. (This the content which will be placed in file with filename)

[Ctrl D] End of input

\$\_ (comes to the shell prompt)

The above command will create a file named os.txt in the current directory. To see the contents of the file.

```
$ cat os.txt[Enter]
```

Welcome to Manipal.

To display a file we can use **cat** command as shown above.

We can use **cat** for displaying more than one file, one after the other by listing the files after **cat**. For e.g.

```
$ cat os.txt lab.txt
```

will display os.txt followed with lab.txt **cp:** To copy the contents of one file to another.

Syntax: **cp** sourcefilename targetfilename [Enter]

This command is also used to copy one or more files to a directory. The syntax of this form of **cp** command is

Syntax : **cp** filename(s) directoryname

If the file **os.txt** in current directory i.e. /home/kumar/pis needs to be copied into /home directory then it will be done as follows.

```
$ cp os.txt /home/ OR $ cp os.txt ../.. mv: This command renames or moves files. It has two distinct function: It renames a file or a directory and it moves a group of files to a different directory.
```

**Syntax: mv** oldfilename newfilename Syntax of another form of this command is

**mv** file(s) directory **mv** doesn't create a copy of the file, it merely renames it. No additional space is consumed on disk for the file after renaming. To rename the file chap01 to man01,

```
$ mv chap01 man01.
```

If the destination file doesn't exist, it will be created. For the above example, **mv** simply replaces the filename in the existing directory with the new name. By default **mv** doesn't prompt for overwriting the destination file if it exist.

The following command moves three files to the progs directory: \$ **mv**  
chap01 chap02 chap03 progs

**mv** can also be used to rename a directory for instance pis to pos:

\$ **mv** pis pos **rm**: This command deletes one or more files. Syntax: **rm** filename

The following command deletes three files

```
$ rm chap01 chap02 chap03[Enter]
```

A file once deleted can be recovered subject to conditions by using additional software. **rm** won't normally remove a directory but it can

remove files from one or more directories. It can remove two chapters from the  
progs directory by using:

```
$ rm progs/chap01 progs/chap02
```

**mkdir**: Directories are created by **mkdir** command. The command is followed by the name of the directories to be created.

Syntax: **mkdir** directoryname

```
$ mkdir data [Enter]
```

This creates a directory named data under the current directory.

```
$ mkdir data dbs doc
```

The above command creates three directories with names data, dbs and doc.

**rmdir** : Directories are removed by **rmdir** command. The command is followed by the name of the directory to be removed. If a directory is not empty, then the directory will not be removed.

Syntax: **rmdir** directoryname

\$ **rmdir** patch [Enter]

The command removes the directory by the name patch.

In Linux every file and directory has access permissions. Access permissions define which users have permission to access a file or directory. Permissions are three types, read, write and execute. Access permissions are defined for user, group and others.

For e.g. If access permission is only read for user, group and others, then it will be r- -r- -r- -

Access permissions can also be represented as a number. This number is in octal system. An access permission represented in numerical octal format is called absolute permission. The absolute permission for the above is

444

If the access permission is read, write for user, read, execute for group and only execute for others then it will be, rw-r-x- -x

The absolute permission for the above is

651

**chmod**: changes the permission specified in the argument and leaves the other permissions unaltered. In this mode the following is the syntax.

Syntax: **chmod** category operation permission filename(s) **chmod** takes as its argument an expression comprising some letters and symbols that completely describe the user category and the type of permission being assigned or removed. The expression contains three components:

User category (user, group, others)

The operation to be performed (assign or remove a permission). The type of permission (read, write and execute)

The abbreviations used for these three components are shown in Table 1.1.

E.g. to assign execute permission to the user of the file xstart;

```
$ chmod u+x xstart
```

```
$ ls -l xstart
```

```
- rwxr--r-- 1 kumar      metal  1980  May 01 20:30 xstart.
```

The command assigns (+) execute (x) permission to the user (u), but other permissions remain unchanged. Now the owner of the file can execute the file but the other categories i.e. group and others still can't. To enable all of them to execute this file:

```
$ chmod ugo+x xstart
```

```
$ ls -l xstart
```

```
- rwxr-xr-x 1 kumar      metal  1980  May 01 20:30 xstart.
```

The string **ugo** combines all the three categories user, group and others. This command accepts multiple filenames in the command line:

```
$ chmod u+x note note1 note3
```

```
$ chmod a-x, go+r xstart; ls -l xstart (Two commands can be run simultaneously with ;)
```

```
- rw-r--rwx 1 kumar      metal  1980  May 01 20:30 xstart.
```

**Table 1.1: Abbreviations Used by chmod**

Category	Operation	Permission
u- User	+ Assigns permission	r- Read permission
g- Group	- Removes permission	w- Write permission
o- Others	= Assigns absolute permission	x- Execute permission
a- All(ugo)		

**Absolute Permissions:**

Sometimes without needing to know what a file's current permissions the need to set all nine permission bits explicitly using **chmod** is done.

Read permission – 4 (Octal 100)

Write permission – 2 (Octal 010)

Execute permission – 1 (Octal 001)

For instance, 6 represents read and write permissions, and 7 represents all permissions as can easily be understood from Table 1.2.

**Table 1.2: Absolute Permissions**

Binary	Octal	Permissions	Significance
000	0	---	No permissions
001	1	--x	Executable only
010	2	-w-	Writable only
011	3	-wx	Writable and executable
100	4	r--	Readable only
101	5	r-x	Readable and executable
110	6	rw-	Readable and writable
111	7	rwX	Readable, writable and executable



**\$ chmod 666 xstart; ls -l xstart**

- rw-rw- rw - 1 kumar metal 1980 May 01 20:30 xstart.

The 6 indicates read and write permissions (4 + 2).

**date:** This displays the current date as maintained in the internal clock run perpetually. **\$ date [Enter]**

**clear:** The screen clears and the prompt and cursor are positioned at the top-left corner.

**\$ clear [Enter] man:** is used to display help file related to a command or system call.

Syntax: **man {command name/system call name}** e.g. **man date man open wc:**

displays a count of lines, words and characters in a file. e.g. **wc os.txt** 1 3 19 os.txt

Syntax: **wc [ -c | -m | -C ] [ -l ] [-w] [file....]** Options: The

following options are supported:

-c Count bytes.

-m Count characters.

-C Same as -m,

-l Count lines

-w Count words delimited by white space characters or new line characters.

If no option is specified the default is -lwc (count lines, words, and bytes).

## Redirection Operators

For any program whether it is developed using C, C++ or Java, by default three streams are available known as input stream, output stream and error stream. In programming languages, to refer to them some symbolic names are used (i.e. they are system defined variables).

The following operators are the redirection operators

### 1. > standard output operator

> is the standard output operator which sends the output of any command into a file.

Syntax: **command > file1** e.g. **ls**

**> file1**

Output of the **ls** command is sent to a file1. First, file file1 is created if not exists otherwise, its content is erased and then output of the command is written.

E.g.: **cat file1 > file2**

Here, file2 get the content of file1.

E.g.: **cat file1 file2 file3 > file4**

This creates the file file4 which gets the content of all the files file1, file2 and file3 in order.

## **2. < standard input operator**

< operator (standard input operator) allows a command to take necessary input from a file.

Syntax: **\$ command < file**

E.g.: **cat<file1**

This displays output of file file1 on the screen.

E.g.: **cat <file1 >file2**

This makes cat command to take input from the file file1 and write its output to the file file2. That is, it works like a **cp** command.

## **3. >> appending operator**

Similarly, >> operator can be used to append standard output of a command to a file.

E.g.: **command>>file1**

This makes, output of the given command to be appended to the file1. If the file1 doesn't exist, it will be created and then standard output is written. **4. << document operator**

There are occasions when the data of your program reads is fixed and fairly limited. The shell uses the << symbols to read data from the same file containing the script. This is referred to as **here document**, signifying that the data is here rather than in a separate file. Any command using standard input can also take input from a here document.

Example.:

```
#!/bin/bash  cat
<<DELIMITER
```

```
hello this
is a here
document
DELIMITER
```

This gives the output:

```
hello this
is a here
document
```

### Shell Concepts

This section will describe some of the features that are common in all of the shells.

**1. Wild-card:** The metacharacters that are used to construct the generalized pattern for matching filenames belong to a category called wild-cards.

List of shell's wild-cards:

**Table 1.3 List of Wild Card**

Wild-card	Matches
*	Any number of characters including none
?	A single or zero character
[ijk]	A single character - either an i, j or k
[x-z]	A single character between x and z
[!ijk]	A single character that is not an i, j or k.
[!x-z]	A single character not between x and z.

**Example:** Consider a directory structure /home/kumar which have the following files:

README

chap01

chap02

chap03

helpdir

progs

Then with the below command the following output would be displayed.

**\$ ls chap\***

chap chap01 chap02 chap03

**\$ ls .\***

.bash\_profile .exrc .netscape .profile

- 2. Pipes:** Standard input and standard output constitute two separate streams that can be individually manipulated by the shell. If so then one command can take input from the other. This is possible with the help of pipes.

Assume if the **ls** command which produces the list of files, one file per line, use redirection to save this output to a file:

**\$ ls > user.txt**

**\$ cat user.txt**

The file shows the list of files.

Now to count the number of files:

**\$ ls | wc -l**

The above command gives the number of files. This is how | (pipe) is used.

There's no restriction on the number of commands to be used in pipe.

- 3. Command substitution:** The shell enables the connecting of two commands in yet another way. While a pipe enables a command to obtain its standard input from the standard output of another command, the shell enables one or more command arguments

to be obtained from the standard output of another command. This feature is called command substitution. **\$ echo The date today is `date`**

The date today is Sat May 6 19:01:56 IST 2019

**\$ echo “There are total `ls | wc -l` files and sub-directory in the current directory**

There are 15 files in the current directory.

- 4. Sequences:** Two separate commands can be written in one line using “;” in sequences.  
**\$ chmod 666 xstart; ls -l xstart**

- 5. Conditional Sequences:** The shell provides two operators that allow conditional execution - the && and ||, which typically have this syntax:

cmd1 && cmd2

cmd1 || cmd2

The && delimits two commands; the command cmd2 is executed only when cmd1 succeeds.

The || operator plays inverse role; the second command cmd2 is executed only when the first command cmd1 fails.

Note: All built-in shell commands returns non-zero if they fail. They return zero on success.

e.g: if there is a program hello.c which displays ‘Hello World’ on compilation and execution. Then the following command in conditional sequences could be used to display the same:

**\$ cc hello.c && ./a.out**

This command displays the output ‘Hello World’ if the compilation of the program succeeds. Similarly in case the compilation fails for the program the following output ‘Error’ could be displayed with the following command:

**\$ cc hello.c || echo ‘Error’**

### **File Filters commands in Linux:**

- 1. head:** To see the top 10 lines of a file - **\$ head <file name>** To see the top 5 lines of a file - **\$ head -5 <file name>**
- 2. tail:** To see last 10 lines of a file - **\$ tail <file name>**  
To see last 20 lines of a file - **\$ tail -20 <file name>**

3. **more:** To see the contents of a file in the form of page views - \$ more <file name>

\$ **more fl.txt**

4. **grep:** To search a pattern of word in a file, **grep** command is used.

Syntax: \$ **grep < word name> < file name>**

\$ **grep hi file\_1**

To search multiple words in a file

\$ **grep -E 'word1|word2|word3' <file name>**

\$ **grep -E 'hi|beyond|good' file\_1**

5. **sort:** This command is used to sort the file.

\$ **sort <file name>**

\$ **sort file\_1**

To sort the files in reverse order

\$ **sort -r <file name>**

To display only files

\$ **ls -l | grep "^-"**

To display only directories

\$ **ls -l | grep "^d"**

### **Lab Exercises:**

1. Write shell commands for the following.

i) To create a directory in your home directory having 2 subdirectories.

ii) In the first subdirectory, create 3 different files with different content in each of them.

iii) Copy the first file from the first subdirectory to the second subdirectory.

iv) Create one more file in the second subdirectory, which has the output of the number of users and number of files.

v) To list all the files which starts with either a or A. vi) To count the number of files in the current directory

vii) To count the number of files in the current directory.

viii) Display the output if the compilation of a program succeeds.

- viii) Count the number of lines in an input file.
2. Execute the following commands in sequence: i) date ii) ls iii) pwd

**Additional Exercises:**

1. Write shell commands for the following.
  - i) To display an error message if the compilation of a program fails.
  - ii) To write a text block into a new file. iii. List all the files.
  - iii) To count the number of users logged on to the system.

**LAB NO.: 2**

**Date:**

## **SHELL SCRIPTING 1**

### **Objectives:**

In this lab, the student will be able to:

- Understand the importance of scripts
- Write and execute shell scripts

The Linux shell is a program that handles interaction between the user and the system. Many of the commands that are typically thought of as making up the Linux system are provided by the shell. Commands can be saved as files called scripts, which can be executed like a program.

### **SHELL PROGRAMS: SCRIPTS**

#### **SYNTAX: scriptname**

NOTE: A file that contains shell commands is called a script. Before a script can be run, it must be given execute permission by using **chmod** utility (chmod +x script). To run the script, only type its name. They are useful for storing commonly used sequences of commands to full-blown programs.

### **VARIABLES**

**Table 2.1 :Parameter Variables**

\$@	an individually quoted list of all the positional parameters
\$#	the number of positional parameters
\$_	the process ID of the last background command
\$0	The name of the shell script.
\$\$	The process ID of the shell script, often used inside a script for generating unique temporary filenames; for example /tmp/tmpfile_\$\$.



\$1, \$2, ...	The parameters given to the script.
\$*	A list of all the parameters, in a single variable, separated by the first character in the environment variable IFS.

### Lab Exercises:

1. Try the following shell commands  

```
$ echo $HOME, $PATH
```

```
$ echo $MAIL
```

```
$ echo $USER, $SHELL, $TERM
```
2. Try the following snippet, which illustrates the difference between local and environment variable:  

```
$ firstname=Rakesh .....local variables
```

```
$ lastname=Sharma
```

```
$ echo $firstname $lastname
```

```
$ export lastname .....make "lastname" an envi var
```

```
$ sh .....start a child shell
```

```
$ echo $firstname $lastname
```

```
$ ^D .....terminate child shell
```

```
$ echo $firstname $lastname
```
3. Try the following snippet, which illustrates the meaning of special local variables:  

```
$ cat >script.sh
```

```
echo the name of this script is $0
```

```
echo the first argument is $1
```

```
echo a list of all the arguments is $*
```

```
echo this script places the date into a temporary file
```

```
echo called $1.$$  
date > $1.$$ # redirect the output of date  
ls $1.$$ # list the file  
rm $1.$$ # remove the file
```

^D

```
$ chmod +x script.sh
```

```
$ ./script.sh Rahul Sachin Kumble
```

**NOTE:** A shell supports two kinds of variables: local and environment variables. Both hold data in a string format. The main difference between them is that when a shell invokes a subshell, the child shell gets a copy of its parent shell's environment variables, but not its local variables. Environment variables are therefore used for transmitting useful information between parent shells and their children. **Few predefined environment variables:**

\$HOME pathname of our home directory

\$PATH list of directories to search for commands

\$MAIL pathname of our mailbox

\$USER our username

\$SHELL pathname of our login shell

\$TERM type of the terminal

**Creating a local variable:** variableName=value

### **Operations:**

- Simple assignment and access
- Testing of a variable for existence
- Reading a variable from standard input
- Making a variable read only
- Exporting a local variable to the environment

**Creating / Assigning a variable**

Syntax: {name=value}

Example: \$ firstName=Anand lastname=Sharma age=35

\$ echo \$firstname \$lastname \$age

\$ name = Anand Sharma

\$ echo \$name

\$ name = "Anand Sharma"

\$ echo \$name Accessing  
variable:

Syntax: \$name / \${name}

Example: \$ verb=sing

\$ echo I like \$verbing

Reading a variable from standard input:

Syntax: read {variable}+

Example: \$ cat > script.sh

echo "Please enter your name:" read name echo

your name is \$name

^D

Read-only variables:

Syntax: readonly {variable}+

Example: \$ password=manipal

\$ echo \$password

\$ readonly password

\$ readonly .....list

\$ password=mangalore

## Running jobs in Background

A multitasking system lets a user do more than one job at a time. Since there can be only one job in foreground, the rest of the jobs have to run in the background. There are two ways of doing this: with the shell's **& operator** and **nohup** command. The latter permits to log out while the jobs are running, but the former doesn't allow that.

```
$ sort -o emp.lst &
```

```
550
```

The shell immediately returns a number the PID of the invoked command (550). The prompt is returned and the shell is ready to accept another command even though the previous command has not been terminated yet. The shell however remains the parent of the background process. Using an **&** many jobs can be run in background as the system load permits.

In the above case, if the shell which has started the background job is terminated, the background job will also be terminated. **nohup** is a command for running a job in background in which case the background job will not be terminated if the shell is close. **nohup** stands for no hang up. e.g.

```
$ nohup sort-o emp.lst &
```

```
586
```

The shell returns the PID too. When the **nohup** command is run it sends the standard output of the command to the file **nohup.out**. Now the user can log out of the system without aborting the command.

## JOB CONTROL

**1. ps:** **ps** is a command for listing processes. Every process in a system will have unique id called process id or PID. This command when used displays the process attributes.

```
$ ps
```

```
PID TTY TIME CMD
```

```
291 console 0:00 bash
```

This command shows the PID, the terminal TTY with which the process is associated, the cumulative processor time that has been consumed since the process has started and the process name (CMD).

2. **kill:** This command sends a signal usually with the intention of killing one or more process. This command is an internal command in most shells. The command uses one or more PIDs as its arguments and by

default sends the SIGTERM(15) signal.  
Thus:

\$ **kill 105** terminates the job having PID 105. The command can take many PIDs at a time to be terminated.

3. **sleep:** This command makes the calling process sleep until the specified number of seconds or a signal arrives which is not ignored. \$ **sleep 2**

### Lab Exercises:

1. Try the following, which illustrates the usage of **ps**:

```
$ (sleep 10; echo done) &  
$ ps
```

2. Try the following, which illustrates the usage of **kill**:

```
$ (sleep 10; echo done) &  
$ kill pid .....pid is the process id of background process
```

3. Try the following, which illustrates the usage of **wait**:

```
$ (sleep 10; echo done 1 ) & $ (sleep 10; echo done 2 )  
&
```

```
$ echo done 3; wait ; echo done 4 ....wait for children
```

**NOTE:** The following two utilities and one built-in command allow the listing controlling the current processes.

**ps:** generates a list of processes and their attributes, including their names, process ID numbers, controlling terminals, and owners **kill:** allows to terminate a process on the basis of its ID number **wait:** allows a shell to wait for one or all of its child processes to terminate

Sample Program: \$ cat>script.sh echo there are \$# command line arguments: \$@

```
^D
```

```
$ script.sh arg1 arg2
```

Example: `#!/bin/sh salutation="Hello" echo $salutation echo "The program $0 is now running" echo "The second parameter was $2" echo "The first parameter was $1" echo "The parameter list was $*" echo "The user's home directory is $HOME" echo "Please enter a new greeting" read salutation echo $salutation echo "The script is now complete"`

`exit 0`

If we save the above shell script as `try.sh`, we get the following output: `$ ./try.sh foo bar baz`

Hello

The program `./try.sh` is now running

The second parameter was bar

The first parameter was foo

The parameter list was foo bar baz

The user's home directory is `/home/rick`

Please enter a new greeting

Sire

Sire

The script is now complete

`$`

### **Lab Exercises:**

Write Shell Scripts to do the following:

1. List all the files under the given input directory, whose extension has only one character
2. Write a shell script that accepts two command line parameters. First parameter indicates the directory and the second parameter indicates a regular expression. The script should display all the files and directories in the directory specified in the first argument matching the format specified in the second argument.
3. Count the number of users logged on to the system. Display the output as Number of users logged into the system.

4. Count only the number of files in the current directory.
5. Write a shell script that takes two sorted numeric files as input and produces a single sorted numeric file without any duplicate contents.
6. Write a shell script that accepts two command line arguments. First argument indicates format of file and the second argument indicates the destination directory. The script should copy all the files as specified in the first argument to the location indicated by the second argument.

Also, try the script where the destination directory name has space in it.

### **Additional Exercises:**

1. Write Shell Scripts to do the following
  - i) To list all the .c files in any given input subdirectory.
  - ii) Write a script to include n different commands.

## **SHELL SCRIPTING 2**

### **Objectives:**

In this lab, the student will be able to:

- Grasp the utility of the various variables in the Linux operating system.
- Understand the different arithmetic and relational operators
- Understand the syntax and working of the various looping and decision statements

The shell is not just a collection of commands but a really good programming language. A lot of tasks could be automated with it, along with this the shell is very good for system administration tasks. Many of the ideas could be easily tried with it thus making it as a very useful tool for simple prototyping and it is very useful for small utilities that perform some relatively simple tasks where efficiency is less important as compared to the ease of configuration, maintenance and portability.

### **COMMENTS**

Comments in shell programming start with # and go until the end of the line.

List variables

Syntax: declare [-ax] [listname]

Example: \$ declare -a teamnames

\$ teamnames[0] = "India" .....assignment

\$ teamnames[1] = "England"

\$ teamnames[2] = "Nepal"

\$ echo "There are \${#teamnames[\*]} teams" ....accessing

\$ echo "They are: \${teamnames [\*]}"

\$ unset teamnames[1] ...delete

### **Aliases**

Allows to define your own commands



Syntax: alias [word[=string]]

Unalias [-a] { word }+

Example: \$ alias dir="ls -aF"

\$ dir

## ARITHMETIC

expr utility is used for arithmetic operations. All of the components of expression must be separated by blanks, and all of the shell meta characters must be escaped by a \.

**Syntax: expr expression**

Example: \$ x=1

\$ x=`expr \$x +1`

\$ echo \$x

\$ x=`expr 2 + 3 \\* 5`

\$echo \$x

\$echo `expr \( 4 \> 5 \)`

\$echo `expr length "cat"`

\$echo `expr substr "donkey" 4 3`

## TEST EXPRESSION

Syntax: test expression

Table 3.1 :Forms of Test Expressions

Test	Meaning
!=	not equal
=	equal
-eq	equal

-gt	greater than
-ge	greater than or equal
-lt	less than
-le	less than or equal
!	logic negation
-a	logical and
-o	logical or
-r file	true if the file exists and is readable
-w file	true if the file exists and is writable
-x file	true if the file exists and is executable
-s file	true if the file exists and its size > 0
-d file	true if the file is a directory
-f file	true if the file is an ordinary file
-t file	true if the file descriptor is associated with a terminal
-n str	true if the length of str is > 0
-z str	true if the length of str is zero

## CONTROL STRUCTURES

### (i) The if conditional

Syntax: if command1

then command2 fi

Example: *echo*

*“enter a  
number:”*

*read number if*

*[ \$number -lt 0*

*] then echo*

*“negative” elif*

*[ \$number -eq*

```

0 ] then echo
    "zero" else
    echo "positive"
fi

```

## (ii) The case conditional

Syntax:

```

case string in pattern1)
    commands1 ;; pattern2)
    commands2 ;; ..... esac

```

case selectively executes statements if string matches a pattern. You can have any number of patterns and statements. Patterns can be literal text or wildcards. You can have multiple patterns separated by the "|" character.

### Example:

```

case $1 in
*.c)
cc $1
;;
*.h | *.sh)
# do nothing
;;

```

The above example performs a compile if the filename ends in .c, does nothing for files ending in .h or .sh. else it writes to stdout that the file is an unknown type. Note that the: character is a NULL command to the shell (similar to a comment field).

```

case $1 in
[AaBbCc])

```

```

option=0

```

```

;; *)
optio
n=1
;; esac
echo
$option

```

In the above example, if the parameter \$1 matches A, B or C (uppercase or lowercase), the shell variable *option* is assigned the value 0, else is assigned the value 1.

### (iii) while: looping

Syntax:

```
while condition is true do
    commands
```

```
done
```

#### **Example 1:**

```
# menu program echo "menu test program" stop=0 while test $stop -eq 0 do cat <<
ENDOFMENU
```

```
1: print the date
```

```
2,3    : print the current working directory
```

```
4: exit
```

```
ENDOFMENU echo echo "your choice ?" read reply echo case $reply in
"1") date
```

```
;;
```

```
"2" | "3") pwd
```

```
;;
```

```
"4") stop =1
```

```
;;
```

```
*) echo "illegal choice"
```

```
;; esac
```

```
done
```

#### **Example 2:**

```
#!/bin/bash X=0
```

```
while [ $X -le 20 ] do
```

```
    echo $X
```

```
    X=$((X+1)) done
```

```
# echo all the command line arguments
```

```
while test $# != 0 do
```

```
echo $1      #The shift command shifts arguments to the left
```

```
shift
```

```
done
```

(iv) **until: Looping**

Syntax: until *command-list1* do *command list2* done

Example:

```
x=1
until [ $x -gt 3
] do echo x =
$x x=`expr $x
+ 1` done
```

(v) **for: Looping**

Syntax: for variable in list do command-list done

Sample Program

```
homedir=`pwd`
for files in
/* do echo
$files done
cd $homedir
```

The above example lists the names of all files under / (the root directory)

### **Lab Exercises:**

Write shell scripts to perform the following

1. Find whether the given number is even or odd.
2. Print the first 'n' odd numbers.
3. Find all the possible quadratic equation roots using case.
4. Find the factorial of a given number.

### **Additional Exercises**

Write Shell scripts to do the following

1. Find whether the given string is palindrome.
2. Find out the sum of the numbers given by user.

## **LINUX SYSTEM CALLS**

### **Objectives:**

In this lab, the student will be able to:

- Learn to create files, open them, read, write and close them
- Learn the need to make system calls and a whole range of library functions to efficiently handle files

Each running program, called a process, has several file descriptors associated with it. When a program starts, it usually has three of these descriptors already opened. These are:

0: Standard input

1: Standard output

2: Standard error

### **SYSTEM CALLS:**

#### **Write System Call:**

The write system call arranges for the first nbytes bytes from buf to be written to the file associated with the file descriptor fildes. It returns the number of bytes written. This may be less than nbytes if there has been an error in the file descriptor or if the underlying device driver is sensitive to block size. If the function returns 0, it means no data was written; if it returns -1, there has been an error in the write call, and the error will be specified in the errno global variable.

Here's the syntax:

```
#include <unistd.h>
```

```
size_t write(int fildes, const void *buf, size_t nbytes);
```

```
#include <unistd.h>
```

```
#include <stdlib.h>
```

```
int main()
```

```
{
```

```
if ((write(1, "Here is some data\n", 18)) != 18)
```

```
write(2, "A write error has occurred on file descriptor 1\n", 46);
```

```
exit(0);
```

```
}
```

This program simply prints a message to the standard output. When a program exits, all open file descriptors are automatically closed, so you don't need to close them explicitly. This won't be the case, however, when you're dealing with buffered output.

```
$ ./simple_write
Here is some data
$
```

A point worth noting again is that write might report that it wrote fewer bytes than you asked it to. This is not necessarily an error. In your programs, you will need to check `errno` to detect errors and call write to write any remaining data.

### **Read System Call:**

The read system call reads up to `nbytes` bytes of data from the file associated with the file descriptor `fildev` and places them in the data area `buf`. It returns the number of data bytes read, which may be less than the number requested. If a read call returns 0, it had nothing to read; it reached the end of the file. Again, an error on the call will cause it to return `-1`.

```
#include <unistd.h>
size_t read(int fildev, void *buf, size_t nbytes);
This program, simple_read.c, copies the first 128 bytes of the standard input to the
standard output. It
copies all of the input if there are fewer than 128 bytes.
#include <unistd.h>
#include <stdlib.h>
int main()
{
char buffer[128];
int nread;
nread = read(0, buffer, 128);
if (nread == -1)
write(2, "A read error has occurred\n", 26);
if ((write(1,buffer,nread)) != nread)
write(2, "A write error has occurred\n",27); exit(0);
}
```

If you run the program, you should see the following:

```
$ echo hello there | ./simple_read
hello there
$ ./simple_read < draft1.txt
```

In the first execution, you create some input for the program using `echo`, which is piped to your program. In the second execution, you redirect input from a file. In this case, you see the first part of the file `draft1.txt` appearing on the standard output.

## Open System Call:

To create a new file descriptor, you need to use the open system call.

```
#include <fcntl.h>
#include <sys/types.h>
#include <sys/stat.h>
int open(const char *path, int oflags);
int open(const char *path, int oflags, mode_t mode);
```

In simple terms, open establishes an access path to a file or device. If successful, it returns a file descriptor that can be used in read, write, and other system calls. The file descriptor is unique and isn't shared by any other processes that may be running. If two programs have a file open at the same time, they maintain distinct file descriptors. If they both write to the file, they will continue to write where they left off. Their data isn't interleaved, but one will overwrite the other. Each keeps its idea of how far into the file (the offset) it has read or written. You can prevent unwanted clashes of this sort by using file locking. The name of the file or device to be opened is passed as a parameter, path; the oflags parameter is used to specify actions to be taken on opening the file. The oflags are specified as a combination of a mandatory file access mode and other optional modes. The open call must specify one of the file access modes shown in the following table:

Mode	Description
O_RDONLY	Open for read-only
O_WRONLY	Open for write-only
O_RDWR	Open for reading and writing

The call may also include a combination (using a bitwise OR) of the following optional modes in the oflags parameter:

**O\_APPEND:** Place written data at the end of the file.

**O\_TRUNC:** Set the length of the file to zero, discarding existing contents.

**O\_CREAT:** Creates the file, if necessary, with permissions given in mode.

**O\_EXCL:** Used with O\_CREAT, ensures that the caller creates the file. The open is atomic; that is, it's performed with just one function call. This protects against two programs creating the file at the same time. If the file already exists, open will fail.

Other possible values for oflags are documented in the open manual page, which you can find in section 2 of the manual pages (use `man 2 open`).

open returns the new file descriptor (always a nonnegative integer) if successful, or -1 if it fails, at which time open also sets the global variable errno to indicate the reason for the failure. We look at errno more closely in a later section. The new file descriptor is always the lowest-numbered unused descriptor, a feature that can be quite useful in some



circumstances. For example, if a program closes its standard output and then calls open again, the file descriptor 1 will be reused and the standard output will have been effectively redirected to a different file or device.

## Initial Permission

When you create a file using the O\_CREAT flag with open, you must use the three-parameter form. mode, the third parameter, is made from a bitwise OR of the flags defined in the header file sys/stat.h. These are:

S\_IRUSR: Read permission, owner

S\_IWUSR: Write permission, owner

S\_IXUSR: Execute permission, owner

S\_IRGRP: Read permission, group

S\_IWGRP: Write permission, group

S\_IXGRP: Execute permission, group

S\_IROTH: Read permission, others

S\_IWOTH: Write permission, others

S\_IXOTH: Execute permission, others

For example,

```
open ("myfile", O_CREAT, S_IRUSR|S_IXOTH);
```

## Unmask system call

The umask is a system variable that encodes a mask for file permissions to be used when a file is created. You can change the variable by executing the umask command to supply a new value. The value is a three-digit octal value. Each digit is the result of ORing values from 1, 2, or 4; the meanings are shown in the following table. The separate digits refer to "user," "group," and "other" permissions, respectively.

Digit	Value	Meaning
1	0	No user permissions are to be disallowed.
	4	User read permission is disallowed.
	2	User write permission is disallowed.
	1	User execute permission is disallowed.
2	0	No group permissions are to be disallowed.
	4	Group read permission is disallowed.
	2	Group write permission is disallowed.
	1	Group execute permission is disallowed.
3	0	No other permissions are to be disallowed.
	4	Other read permission is disallowed.
	2	Other write permission is disallowed.
	1	Other execute permission is disallowed

## Close System Call

You use close to terminate the association between a file descriptor, fildes, and its file. The file descriptor becomes available for reuse. It returns 0 if successful and -1 on error.

```
#include <unistd.h>
int close(int fildes);
```

Note that it can be important to check the return result from close. Some file systems, particularly networked ones, may not report an error writing to a file until the file is closed, because data may not have been confirmed as written when writes are performed.

## A File Copy Program

```
#include <unistd.h>
#include <sys/stat.h>
#include <fcntl.h>
#include <stdlib.h>
int main()
{
    char c;
    int in, out;
    in = open("file.in", O_RDONLY);
    out = open("file.out", O_WRONLY|O_CREAT, S_IRUSR|S_IWUSR);
    while(read(in,&c,1) == 1)
        write(out,&c,1);
    exit(0);
}
```

WITH fcntl.h how the system calls can manipulate the files.

## I/O SYSTEM CALLS

I/O through system calls is simpler and operates at a lower level than making calls to the C file-I/O library.

There are seven fundamental file-I/O system calls:

- creat() Create a file for reading or writing.
- open() Open a file for reading or writing.
- close() Close a file after reading or writing.
- unlink() Delete a file.
- write() Write bytes to file.
- read() Read bytes from file.

## The creat() System Call

The "creat()" system call creates a file. It has the syntax: `int fp; /* fp is the file descriptor variable */`

`fp = creat( <filename>, <protection bits> );`

Ex: `fp=creat("students.dat",RD_WR);`

This system call returns an integer, called a "file descriptor", which is a number that identifies the file generated by "creat()". This number is used by other system calls in the program to access the file. Should the "creat()" call encounter an error, it will return a file descriptor value of -1.

The "filename" parameter gives the desired filename for the new file.

The "permission bits" give the "access rights" to the file. A file has three "permissions" associated with it:

Write permission - Allows data to be written to the file.

Read permission - Allows data to be read from the file.

Execute permission - Designates that the file is a program that can be run.

These permissions can be set for three different levels:

User level: Permissions apply to individual user.

Group level: Permissions apply to members of user's defined "group".

System level: Permissions apply to everyone on the system

## The open() System Call

The "open()" system call opens an existing file for reading or writing. It has the syntax:

`<file descriptor variable> = open( <filename>, <access mode> );`

The "open()" call is similar to the "creat()" call in that it returns a file descriptor for the given file, and returns a file descriptor of -1 if it encounters an error. However, the second

parameter is an "access mode", not a permission code. There are three modes (defined in the "fcntl.h" header file):

O_RDONLY	Open for reading only.
O_WRONLY	Open for writing only.
O_RDWR	Open for reading and writing

For example, to open "data" for writing, assuming that the file had been created by another program, the following statements would be used:

```
int fd;
```

```
fd = open( "students.dat", O_WRONLY );
```

A few additional comments before proceeding:

A "creat()" call implies an "open()". There is no need to "creat()" a file and then "open()" it.

### **The close() System Call**

The "close()" system call is very simple. All it does is "close()" an open file when there is no further need to access it. The "close()" system call has the syntax:

```
close( <file descriptor> );
```

The "close()" call returns a value of 0 if it succeeds, and returns -1 if it encounters an error.

### **The write() System Call**

The "write()" system call writes data to an open file. It has the syntax:

```
write( <file descriptor>, <buffer>, <buffer length> );
```

The file descriptor is returned by a "creat()" or "open()" system call. The "buffer" is a pointer to a variable or an array that contains the data; and the "buffer length" gives the number of bytes to be written into the file.

While different data types may have different byte lengths on different systems, the "sizeof()" statement can be used to provide the proper buffer length in bytes. A "write()" call could be specified as follows:

```
float array[10];  
write( fd, array, sizeof( array ) );
```

The "write()" function returns the number of bytes it writes. It will return -1 on an error.

## The read() Sytem Call

The "read()" system call reads data from an open file. Its syntax is the same as that of the "write()" call:

```
read( <file descriptor>, <buffer>, <buffer length> );
```

The "read()" function returns the number of bytes it returns. At the end of the file, it returns 0 or returns -1 on error.

**lseek:** The lseek system call sets the read/write pointer of a file descriptor, fildes; that is, we can use it to set wherein the file the next read or write will occur. We can set the pointer to an absolute location in the file or a position relative to the current position or the end of file.

```
#include <unistd.h>
```

```
#include <sys/types.h>
```

```
off_t lseek(int fildes, off_t offset, int whence);
```

The offset parameter is used to specify the position, and the whence parameter specifies how the offset is used. whence can be one of the following:

SEEK\_SET: offset is an absolute position

SEEK\_CUR: offset is relative to the current position

SEEK\_END: offset is relative to the end of the file

**lseek** returns the offset measured in bytes from the beginning of the file that the file pointer is set to or -1 on failure. The type off\_t, used for the offset in seek operations, is an implementation-dependent type defined in sys/types.h.

### Errors:

EACCES Permission denied.

EMFILE Too many file descriptors in use by the process.

ENFILE Too many files are currently open in the system.

ENOENT Directory does not exist, or the name is an empty string.

ENOMEM Insufficient memory to complete the operation.

### **Lab Exercises:**

2. Write a program to print the lines of a file that contain a word given as the program argument (a simple version of grep UNIX utility).
3. Write a program to list the files given as arguments, stopping every 20 lines until a key is hit. (a simple version of more UNIX utility)
4. Demonstrate the use of different conversion specifiers and resulting output to allow the items to be printed.
5. Write a program to copy character-by character copy is accomplished using calls to the functions referenced in stdio.h

### **Additional Exercises:**

1. Write a program that shows the user all his/her C source programs and then prompts interactively as to whether others should be granted read permission; if affirmative such permission should be granted.
2. Use lseek() to copy different parts (initial, middle and last) of the file to others. (For lseek() refer to man pages)

## PROCESSES AND SIGNALS

### Objectives:

In this lab, the student will be able to

- Learn how the operating system manages processes
- Learn how the system handles processes, processes send, and receive messages
- Understand the concept of orphan and zombie processes

### System Calls Related to Processes

#### **getpid()**

This function returns the process identifiers of the calling process.

```
#include <sys/types.h>
```

```
#include <unistd.h>
```

```
pid_t getpid(void); // this function returns the process identifier (PID)
```

```
pid_t getppid(void); // this function returns the parent process identifier (PPID)
```

#### **fork()**

A new process is created by calling fork. This system call duplicates the current process, creating a new entry in the process table with many of the same attributes as the current process. The new process is almost identical to the original, executing the same code but with its own data space, environment, and file descriptors. Combined with the **exec** functions, the **fork** is all we need to create new processes.

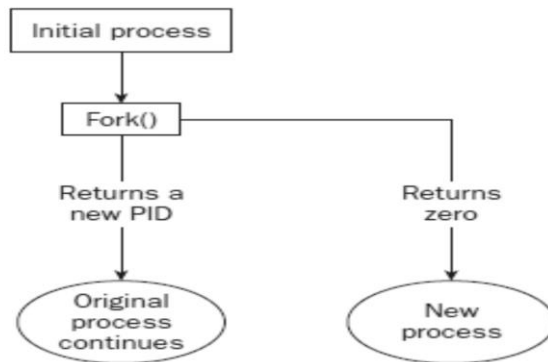
```
#include <sys/types.h>
```

```
#include <unistd.h>
```

```
pid_t fork(void);
```

The return value of fork() is pid\_t (defined in the header file sys/types.h). As seen in Fig. 4.1, the call to the fork in the parent process returns the PID of the new child process. The new process continues to execute just like the parent process, with the exception that in the child process, the PID returned is 0. The parent and child process can be determined by using the PID returned from the fork() function. To the parent, the fork() returns the

PID of the child, whereas to the child the PID returned is zero. This is shown in the following Fig. 5.1.



**Figure 5.1: Fork system call**

In Linux in case of any error observed in calling the system functions, then a special variable called `errno` will contain the error number. To use `errno` a header file named `errno.h` has to be included in the program. If the fork fails, it returns -1. This is commonly due to a limit on the number of child processes that a parent may have (`CHILD_MAX`), in which case `errno` will be set to `EAGAIN`. If there is not enough space for an entry in the process table, or not enough virtual memory, the `errno` variable will be set to `ENOMEM`.

A typical code snippet using fork is

```
pid_t new_pid;
new_pid = fork();
switch(new_pid) {
case -1 : /* Error */
break;
case 0 : /* We are child */
break;
default : /* We are parent */
break;
}
```

### **Sample Program on fork1.c**

```
#include <sys/types.h>
#include <unistd.h>
#include <stdio.h>
```



```

int main()
{
    pid_t pid;
    char *message;
    int n;
    printf("fork program starting\n");
    pid = fork();
    switch(pid)
    {
        case -1:
            perror("fork failed");
            exit(1);
        case 0:
            message = "This is the child";
            n = 5;
            break;
        default:
            message = "This is the parent";
            n = 3;
            break;
    }
    for(; n > 0; n--) {
        puts(message);
        sleep(1);
    }
    exit(0);
}

```

This program runs as two processes. A child process is created and prints a message five times. The original process (the parent) prints a message only three times. The parent process finishes before the child has printed all of its messages, so the next shell prompt appears mixed in with the output.

```
$ ./fork1
```

```
fork program starting
```

```
This is the parent
```

```
This is the child
```

This is the parent

This is the child

This is the parent

This is the child

This is the child

This is the child

When fork is called, this program divides into two separate processes. The parent process is identified by a nonzero return from fork and is used to set several messages to print, each separated by one second.

### **The wait() System Call**

A parent process usually needs to synchronize its actions by waiting until the child process has either stopped or terminated its actions. The wait() system call allows the parent process to suspend its activities until one of these actions has occurred. The wait() system call accepts a single argument, which is a pointer to an integer and returns a value defined as type pid\_t. If the calling process does not have any child associated with it, the wait will return immediately with a value of -1. If any child processes are still active, the calling process will suspend its activity until a child process terminates.

Example of wait():

```
#include <sys/types.h>
```

```
#include <sys/wait.h>
```

```
void main()
```

```
{
```

```
    int status;
```

```
    pid_t pid;
```

```
    pid = fork();
```

```
    if(pid == -1)
```

```
        printf("\nERROR child not created");
```

```
    else if (pid == 0) /* child process */
```

```
    {
```

```
        printf("\n I'm the child!");
```

```
        exit(0);
```

```
    }
```

```
    else /* parent process */
```

```
    {
```

```

        wait(&status);
        printf("\n I'm the parent!")
        printf("\n Child returned: %d\n", status)
    }
}

```

A few notes on this program:

wait(&status) causes the parent to sleep until the child process has finished execution. The exit status of the child is returned to the parent.

### **The exit() System Call**

This system call is used to terminate the currently running process. A value of zero is passed to indicate that the execution of the process was successful. A non-zero value is passed if the execution of the process was unsuccessful. All shell commands are written in C including grep. grep will return 0 through exit if the command is successfully run (grep could find a pattern in the file). If grep fails to find a pattern in a file, then it will call exit() with a non-zero value. This applies to all commands.

### **The exec() System Call**

The exec function will execute a specified program passed as an argument to it, in the same process (Fig. 5.2). The exec() will not create a new process. As a new process is not created, the process ID (PID) does not change across an execution, but the data and code of the calling process are replaced by those of the new process.

fork() is the name of the system call that the parent process uses to "divide" itself ("fork") into two identical processes. After calling fork(), the created child process is an exact copy of the parent - which would probably be of limited use - so it replaces itself with another process using the system call exec().


### **The versions of exec are:**

- execl
- execv
- execl
- execve
- execlp

- `execvp`

**The naming convention:** `exec*`

- 'l' indicates a list arrangement (a series of null-terminated arguments)
- 'v' indicates the array or vector arrangement (like the `argv` structure).
- 'e' indicates the programmer will construct (in the array/vector format) and pass their environment variable list
- 'p' indicates the current `PATH` string should be used when the system searches for executable files.



**Figure 5.2: `exec()` system call**

The parent process can either continue execution or wait for the child process to complete. If the parent chooses to wait for the child to die, then the parent will receive the exit code of the program that the child executed. If a parent does not wait for the child, and the child terminates before the parent, then the child is called a **zombie** process. If a parent terminates before the child process then the child is attached to a process called `init` (whose `PID` is 1). In this case, whenever the child does not have a parent then the child is called the **orphan** process.

**Sample Program:**

C program forking a separate process.

```

#include<sys/types.h>
#include<stdio.h>
#include<unistd.h>
int main()

{
    pid_t pid;
    /* fork another process */

```

```

pid = fork();
if (pid < 0) { /* error occurred */
    fprintf(stderr, "Fork Failed");
    exit(-1);
}
else if (pid == 0) { /* child process */
    execlp("/bin/ls", "ls", NULL);
}
else { /* parent process */
    /* parent will wait for the child to complete */
    wait (NULL);
    printf ("Child Complete");
    exit(0);
}
}

```

**execl:** is used with a list comprising the command name and its arguments:

```
int execl(const char *path, const char *arg0, ...../*, (char *) 0 */);
```

This is used when the number of arguments is known in advance. The first argument is the pathname which could be absolute or a relative pathname, The arguments to the command to run are represented as separate arguments beginning with the name of the command (\*arg0). The ellipsis representation in the syntax (.../\*) points to the varying number of arguments.

**Example:** How to use execl to run the wc -l command with the filename foo as argument:

```
execl ("/bin/wc", "wc", "-l", "foo", (char *) 0);
```

execl doesn't use PATH to locate wc so pathname is specified as the first argument.

**execv:** needs an array to work with.

```
int execv(const char *path, char *const argv[ ]);
```

Here path represents the pathname of the command to run. The second argument represents an array of pointers to char. The array is populated by addresses that point to strings representing the command name and its arguments, in the form they are passes to the main function of the program to be executed. In this case, also the last element of the

argv[ ] must be a null pointer.

Here the following program uses execv program to run grep command with two options to look up the author's name in /etc/passwd. The array \*cmdargs[ ] are populated with the strings comprising the command line to be executed by execv. The first argument is the pathname of the command:

```
#include<stdio.h>
int main(int argc, char **argv){
char *cmdargs[ ] = {"grep", "-l", "-n", "SUMIT", "/etc/passwd", NULL};
execv("/bin/grep", cmdargs);
printf ("execv error\n");
}
```

### **Drawbacks:**

Need to know the location of the command file since neither execl nor execv will use PATH to locate it. The command name is specified twice - as the first two arguments. These calls can't be used to run a shell script but only binary executable. The program has to be invoked every time there is a need to run a command.

execlp and execvp: requires the pathname of the command to be located. They behave exactly like their other counterparts but overcomes two of the four limitations discussed above. First the first argument need not be a pathname it can be a command name. Second these functions can also run a shell script.

```
int execlp(const char *file, const char *arg0, .../*, (char *) 0 */);
int execvp(const char *file, char *const argv[ ]);

execlp ("wc", "wc", "-l", "foo", (char *) 0);
```

execle and execve: All of the previous four exec calls silently pass the environment of the current process to the executed process by making available the environ[ ] variable to the overlaid process. Sometime there may be a need to provide a different environment to the new program - a restricted shell for instance. In that case, these functions are used.

```
int execle(const char *path, const char *arg0, ... /*, (char *) 0, char * const envp[ ] */);
int execve(const char *path, char * const argv[ ], char *const envp[ ]);
```

These functions unlike the others use an additional argument to pass a pointer to an array of environment strings of the form variable = value to the program. It's only this

environment that is available in the executed process, not the one stored in `envp[ ]`.

The following program (assume `fork2.c`) is the same as `fork1.c`, except that the number of messages printed by the child and parent processes is reversed. Here are the relevant lines of code:

```
switch(pid)
{
    case -1:
        perror("fork failed");
        exit(1);

    case 0:
        message = "This is the child";
        n = 3;
        break;

    default:
        message = "This is the parent";
        n = 5;
        break;
}
```

When the preceding program is run with `./fork2 &` and then call the `ps` program after the child has finished but before the parent has finished, a line such as this. (Some systems may say `<zombie>` rather than `<defunct>`) is seen.

### **Lab Exercises:**

1. Write a C program to block a parent process until the child completes using a `wait` system call.
2. Write a C program to load the binary executable of the previous program in a child process using the `exec` system call.
3. Write a program to create a child process. Display the process IDs of the process, parent and child (if any) in both the parent and child processes.

4. Create a zombie (defunct) child process (a child with `exit()` call, but no corresponding `wait()` in the sleeping parent) and allow the `init` process to adopt it (after parent terminates). Run the process as a background process and run the “`ps`” command.

### **Additional Exercises**

1. Create an orphan process (parent dies before child – adopted by “`init`” process) and display the PID of the parent of the child before and after it becomes orphan. Use `sleep(n)` in the child to delay the termination.
2. Modify the program in the previous question to include `wait (&status)` in the parent and to display the exit return code (leftmost byte of `status`) of the child.



## IPC-1 : PIPE, FIFO


In this lab, the student will be able to:

- Gain knowledge as to how Interprocess Communication (IPC) happens between two processes
- Execute programs for IPC using the different methods of pipe and fifo

Inter-Process Communication (IPC), is the mechanism whereby one process can communicate with another process, i.e exchange of data. IPC in Linux can be implemented by using a pipe, shared memory and message queue.

### Pipe

- Pipes are unidirectional byte streams that connect the standard output from one process into the standard input of another process. A pipe is created using the system call `pipe` that returns a pair of file descriptors.



**Figure 6.1. Working of pipe**

- As shown in Figure 6.1 and 6.2 call to the `pipe ()` function which returns an array of file descriptors `fd[0]` and `fd [1]`. `fd [1]` connects to the write end of the pipe, and `fd[0]` connects to the read end of the pipe. Anything can be written to the pipe, and read from the other end in the order it came in.
- A pipe is one-directional providing one-way flow of data and it is created by the `pipe()` system call.

```
int pipe ( int *filedes ) ;
```

- An array of two file descriptors are returned- `fd[0]` which is open for reading, and `fd[1]` which is open for writing. It can be used only between parent and child processes.

PROTOTYPE: int pipe( int fd[2] );

RETURNS: 0 on success

-1 on error: errno = EMFILE (no free descriptors)

EMFILE (system file table is full)

EFAULT (fd array is not valid)

fd[0] is set up for reading, fd[1] is set up for writing. i.e., the first integer in the array (element 0) is set up and opened for reading, while the second integer (element 1) is set up and opened for writing.

```
#include <stdlib.h>
```

```
#include <stdio.h>    /* for printf */
```

```
#include <string.h>    /* for strlen */
```

```
int main(int argc, char **argv)
```

```
{
```

```
    int n;
```

```
    int fd[2];
```

```
    char buf[1025];
```

```
    char *data = "hello... this is sample data";
```

```
    pipe(fd);
```

```
    write(fd[1], data, strlen(data));
```

```
    if ((n = read(fd[0], buf, 1024)) >= 0) {
```

```
        buf[n] = 0;    /* terminate the string */
```

```
        printf("read %d bytes from the pipe: \"%s\"\n", n, buf);
```


```
    }
```

```
    else
```

```
        perror("read");
```

```
    exit(0);
```

```
}
```



**Fig. 6.2: Working of pipe in a single process which is immediately after fork()**

- First, a process creates a pipe and then forks to create a copy of itself.
- The parent process closes the read end of the pipe.
- The child process closes the write end of the pipe.
- The fork system call creates a copy of the process that was executing.
- The process which executes the fork is called the parent process and the new process which is created is called the child process.

```
#include <sys/wait.h>
#include <assert.h>
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <string.h>

int main(int argc, char *argv[])
{
    int pfd[2];
    pid_t cpid;
    char buf;
    assert(argc == 2);
    if (pipe(pfd) == -1) { perror("pipe");
        exit(EXIT_FAILURE); }
```

```

cpid = fork();
if (cpid == -1) { perror("fork");
exit(EXIT_FAILURE); }

if (cpid == 0) { /* Child reads from pipe */
    close(pfd[1]); /* Close unused write end */
    while (read(pfd[0], &buf, 1) > 0)
        write(STDOUT_FILENO, &buf, 1);
    write(STDOUT_FILENO, "\n", 1);
    close(pfd[0]);
    exit(EXIT_SUCCESS);

} else { /* Parent writes argv[1] to pipe */
    close(pfd[0]); /* Close unused read end */
    write(pfd[1], argv[1], strlen(argv[1]));
    close(pfd[1]); /* Reader will see EOF */
    wait(NULL); /* Wait for child */
    exit(EXIT_SUCCESS);
}
}

```

## Named Pipes: FIFOs

Pipes can share data between related processes, i.e. processes that have been started from a common ancestor process. We can use named pipe or FIFOs to overcome this. A named pipe is a special type of file that exists as a name in the file system but behaves like the unnamed pipes we have discussed already. We can create named pipes from the command line using

```
$ mkfifo filename
```

From inside a program, we can use

```

#include <sys/types.h>
#include <sys/stat.h>

```

```
int mkfifo(const char *filename, mode_t mode);
```

```
#include <unistd.h>
#include <stdlib.h>
#include <stdio.h>
#include <sys/types.h>
#include <sys/stat.h>
int main()
{
    int res = mkfifo("/tmp/my_fifo", 0777);
    if (res == 0) printf("FIFO created\n");
    exit(EXIT_SUCCESS);
}
```

We can look for the pipe with

```
$ ls -lF /tmp/my_fifo
```

```
prwxr-xr-x 1 rick users 0 July 10 14:55 /tmp/my_fifo|
```

Notice that the first character of output is a p, indicating a pipe. The | symbol at the end is added by the ls command's -F option and also indicates a pipe. We can remove the FIFO just like a conventional file by using the rm command, or from within a program by using the unlink system call.

### **Producer-Consumer Problem (PCP):**

- The producer process produces information that is consumed by a consumer process. To allow producer and consumer processes to run concurrently, we must have available a buffer of items that can be filled by the producer and emptied by the consumer. Two types of buffers can be used.
  - *unbounded-buffer* places no practical limit on the size of the buffer.
  - *bounded-buffer* assumes that there is fixed buffer size.
- For bounded-buffer PCP basic synchronization requirement is:
  - Producer should not write into a full buffer (i.e. producer must wait if the buffer is full)
  - Consumer should not read from an empty buffer (i.e. consumer must wait if the buffer is empty)
  - All data written by the producer must be read exactly once by the consumer

Following is a program for Producer-Consumer problem using named pipes.

```
//producer.c

#include <unistd.h>
#include <stdlib.h>
#include <stdio.h>
#include <string.h>
#include <fcntl.h>
#include <limits.h>
#include <sys/types.h>
#include <sys/stat.h>
#define FIFO_NAME "/tmp/my_fifo"
#define BUFFER_SIZE PIPE_BUF
#define TEN_MEG (1024 * 1024 * 10)

int main()
{
    int pipe_fd;
    int res;
    int open_mode = O_WRONLY;
    int bytes_sent = 0;
    char buffer[BUFFER_SIZE + 1];
    if (access(FIFO_NAME, F_OK) == -1) {
        res = mkfifo(FIFO_NAME, 0777);
        if (res != 0) {
            fprintf(stderr, "Could not create fifo %s\n", FIFO_NAME);
            exit(EXIT_FAILURE);
        }
    }
    printf("Process %d opening FIFO O_WRONLY\n", getpid());
    pipe_fd = open(FIFO_NAME, open_mode);
    printf("Process %d result %d\n", getpid(), pipe_fd);
    if (pipe_fd != -1) {
        while(bytes_sent < TEN_MEG) {
            res = write(pipe_fd, buffer, BUFFER_SIZE);
            if (res == -1) {
                fprintf(stderr, "Write error on pipe\n");
                exit(EXIT_FAILURE);
            }
        }
    }
}
```

```

        bytes_sent += res;
    }
    (void)close(pipe_fd);
}
else {
    exit(EXIT_FAILURE);
}

printf("Process %d finished\n", getpid());
exit(EXIT_SUCCESS);
}

```

//consumer.c

```

#include <unistd.h>
#include <stdlib.h>
#include <stdio.h>
#include <string.h>
#include <fcntl.h>
#include <limits.h>
#include <sys/types.h>
#include <sys/stat.h>
#define FIFO_NAME "/tmp/my_fifo"
#define BUFFER_SIZE PIPE_BUF
int main()
{
    int pipe_fd;
    int res;
    int open_mode = O_RDONLY;
    char buffer[BUFFER_SIZE + 1];
    int bytes_read = 0;
    memset(buffer, '\0', sizeof(buffer));
    printf("Process %d opening FIFO O_RDONLY\n", getpid());
    pipe_fd = open(FIFO_NAME, open_mode);
    printf("Process %d result %d\n", getpid(), pipe_fd);
    if (pipe_fd != -1) {

```

```

        do {
            res = read(pipe_fd, buffer, BUFFER_SIZE);
            bytes_read += res;
        } while (res > 0);
        (void)close(pipe_fd);
    }
    else {
        exit(EXIT_FAILURE);
    }
    printf("Process %d finished, %d bytes read\n", getpid(), bytes_read);
    exit(EXIT_SUCCESS);
}

```

### **The Readers-Writers Problem:**

- Concurrent processes share a file, record, or other resources
- Some may read-only (readers), some may both read and write (writers)
- Two concurrent reads have no adverse effects
- Problems if
  - concurrent reads and writes
  - multiple writes

### **Two Variations**

- First Readers-Writers problem: No reader be kept waiting unless a writer has already obtained exclusive write permissions (Readers have high priority)
- Second Readers-Writers problem: If a writer is waiting/ready, no new readers may start reading (Writers have high priority)

### **Lab Exercises:**

1. Write a producer and consumer program in C using the FIFO queue. The producer should write a set of 4 integers into the FIFO queue and the consumer should display the 4 integers.
2. Demonstrate creation, writing to, and reading from a pipe.
3. Write a C program to implement one side of FIFO.
4. Write two programs, producer.c implementing a producer and consumer.c implementing a consumer, that do the following:  
 Your product will sit on a shelf: that is an integer - a count of the items "on the shelf". This integer may never drop below 0 or rise above 5.



Your producer sets the value of the count to 5. It is the producer program's responsibility to stock product on the shelf, but not overstocked. The producer may add one item to the shelf at a time, and must report to STDOUT every time another item is added as well as the current shelf count.

Your consumer will remove one item from the shelf at a time, provided the item count has not dropped below zero. The consumer will decrement the counter and report the new value to STDOUT. Have your consumer report each trip to the shelf, in which there are no items.

### **Additional Exercises:**

1. Demonstrate creation of a process that writes through a pipe while the parent process reads from it.
2. Demonstrate second readers-writers problem using FIFO

## **IPC – 2 : MESSAGE QUEUE, SHARED MEMORY**

### **Objectives:**

In this lab, the student will be able to:

- Gain knowledge as to how IPC (Interprocess Communication) happens between two processes
- Execute programs for IPC using the different methods of message queues and shared memory

### **Message Queues**

- It is an IPC facility. Message queues are similar to named pipes without the opening and closing of pipe. It provides an easy and efficient way of passing information or data between two unrelated processes.
- The advantages of message queues over named pipes are, it removes a few difficulties that exist during the synchronization, the opening, and closing of named pipes.
- A message queue is a linked list of messages stored within the kernel. A message queue is identified by a unique identifier. Every message has a positive long integer type field, a non-negative length, and the actual data bytes. The messages need not be fetched on FCFS basis. It could be based on the type field.

### **Creating a Message Queue**

- To use a message queue, it has to be created first. The msgget() system call is used for that. This system call accepts two parameters - a queue key and flags.
- IPC\_PRIVATE - use to create a private message queue. A positive integer - used to create or access a publicly accessible message queue.

The message queue function definitions are

```
#include <sys/msg.h>
```

```
int msgctl(int msqid, int cmd, struct msqid_ds *buf);
```

```
int msgget(key_t key, int msgflg);
```

```
int msgrcv(int msqid, void *msg_ptr, size_t msg_sz, long int msgtype, int msgflg);
```

```
int msgsnd(int msqid, const void *msg_ptr, size_t msg_sz, int msgflg);
```

### **msgget**

We create and access a message queue using the msgget function:

**int msgget(key\_t key, int msgflg);**

The program must provide a key-value that, as with other IPC facilities, names a particular message queue. The special value `IPC_PRIVATE` creates a private queue, which in theory is accessible only by the current process. The second parameter, `msgflg`, consists of nine permission flags. A special bit defined by `IPC_CREAT` must be bitwise ORed with the permissions to create a new message queue. It's not an error to set the `IPC_CREAT` flag and give the key of an existing message queue. The `IPC_CREAT` flag is silently ignored if the message queue already exists.

The `msgget` function returns a positive number, the queue identifier, on success or `-1` on failure.

### **msgsnd**

The `msgsnd` function allows us to add a message to a message queue:

**int msgsnd(int msqid, const void \*msg\_ptr, size\_t msg\_sz, int msgflg);**

The structure of the message is constrained in two ways. First, it must be smaller than the system limit, and second, it must start with a long int, which will be used as a message type in the receive function. When you're using messages, it's best to define your message structure something like this:

```
struct my_message {  
    long int message_type;  
    /* The data you wish to transfer */  
}
```

The first parameter, `msqid`, is the message queue identifier returned from a `msgget` function. The second parameter, `msg_ptr`, is a pointer to the message to be sent, which must start with a long int type as described previously. The third parameter, `msg_sz`, is the size of the message pointed to by `msg_ptr`. This size must not include the long int message type. The fourth parameter, `msgflg`, controls what happens if either the current message queue is full or the system-wide limit on queued messages has been reached. If `msgflg` has the `IPC_NOWAIT` flag set, the function will return immediately without sending the message and the return value will be `-1`. If the `msgflg` has the `IPC_NOWAIT` flag clear, the sending process will be suspended, waiting for space to become available in the queue. On success, the function returns `0`, on failure `-1`. If the call is successful, a copy of the message data has been taken and placed on the message queue.

### **msgrcv**

The `msgrcv` function retrieves messages from a message queue:

```
int msgrcv(int msqid, void *msg_ptr, size_t msg_sz, long int msgtype, int msgflg);
```

The first parameter, `msqid`, is the message queue identifier returned from a `msgget` function. The second parameter, `msg_ptr`, is a pointer to the message to be received, which must start with a long int type as described above in the `msgsnd` function. The third parameter, `msg_sz`, is the size of the message pointed to by `msg_ptr`, not including the long int message type. The fourth parameter, `msgtype`, is a long int, which allows a simple form of reception priority to be implemented. If `msgtype` has the value 0, the first available message in the queue is retrieved. If it's greater than zero, the first message with the same message type is retrieved. If it's less than zero, the first message that has a type the same as or less than the absolute value of `msgtype` is retrieved. This sounds more complicated than it is in practice. If you simply want to retrieve messages in the order in which they were sent, set `msgtype` to 0. If you want to retrieve only messages with a specific message type, set `msgtype` equal to that value. If you want to receive messages with a type of `n` or smaller, set `msgtype` to `-n`. The fifth parameter, `msgflg`, controls what happens when no message of the appropriate type is waiting to be received. If the `IPC_NOWAIT` flag in `msgflg` is set, the call will return immediately with a return value of `-1`. If the `IPC_NOWAIT` flag of `msgflg` is clear, the process will be suspended, waiting for an appropriate type of message to arrive. On success, `msgrcv` returns the number of bytes placed in the receive buffer, the message is copied into the user-allocated buffer pointed to by `msg_ptr`, and the data is deleted from the message queue. It returns `-1` on error.

## **msgctl**

The final message queue function is `msgctl`.

```
int msgctl(int msqid, int command, struct msqid_ds *buf);
```

The `msqid_ds` structure has at least the following members:

```
struct msqid_ds {  
    uid_t msg_perm.uid;  
    uid_t msg_perm.gid  
    mode_t msg_perm.mode;  
}
```

The first parameter, `msqid`, is the identifier returned from `msgget`. The second parameter, `command`, is the action to take. It can take three values:

**Table 7.1: Command Description**

Command	Description
IPC_STAT	Sets the data in the <code>msqid_ds</code> structure to reflect the values associated

	with the message queue.
IPC_SET	If the process has permission to do so, this sets the values associated with the message queue to those provided in the msqid_ds data structure.
IPC_RMID	Deletes the message queue.

0 is returned on success, -1 on failure. If a message queue is deleted while a process is waiting in an msgsnd or msgrcv function, the send or receive function will fail.

### Receiver program:

```
#include <stdlib.h>
#include <stdio.h>
#include <string.h>
#include <errno.h>
#include <unistd.h>
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/msg.h>
```

```
struct my_msg_st {
    long int my_msg_type;
    char some_text[BUFSIZ];
};
```

```
int main()
{
    int running = 1;
    int msgid;
    struct my_msg_st some_data;
    long int msg_to_receive = 0;
    msgid = msgget((key_t)1234, 0666 | IPC_CREAT);
    if (msgid == -1) {
        fprintf(stderr, "msgget failed with error: %d\n", errno);
        exit(EXIT_FAILURE);
    }
    while(running) {
        if (msgrcv(msgid, (void *)&some_data, BUFSIZ,
```

```

        msg_to_receive, 0) == -1) {
            fprintf(stderr, "msgrcv failed with error: %d\n", errno);
            exit(EXIT_FAILURE);
        }
        printf("You wrote: %s", some_data.some_text);
        if (strncmp(some_data.some_text, "end", 3) == 0) {
            running = 0;
        }
    }
    if (msgctl(msgid, IPC_RMID, 0) == -1) {
        fprintf(stderr, "msgctl(IPC_RMID) failed\n");
        exit(EXIT_FAILURE);
    }
    exit(EXIT_SUCCESS);
}

```

### **Sender Program:**

```

#include <stdlib.h>
#include <stdio.h>
#include <string.h>
#include <errno.h>
#include <unistd.h>
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/msg.h>

```

```

#define MAX_TEXT 512

```

```

struct my_msg_st {
    long int my_msg_type;
    char some_text[MAX_TEXT];
};

```

```

int main()
{
    int running = 1;
    struct my_msg_st some_data;

```

```

int msgid;
char buffer[BUFSIZ];
msgid = msgget((key_t)1234, 0666 | IPC_CREAT);
if (msgid == -1) {
    fprintf(stderr, "msgget failed with error: %d\n", errno);
    exit(EXIT_FAILURE);
}
while(running) {
    printf("Enter some text:");
    fgets(buffer, BUFSIZ, stdin);
    some_data.my_msg_type = 1;
    strcpy(some_data.some_text, buffer);
    if (msgsnd(msgid, (void *)&some_data, MAX_TEXT, 0) == -1) {
        fprintf(stderr, "msgsnd failed\n");
        exit(EXIT_FAILURE);
    }
    if (strncmp(buffer, "end", 3) == 0) {
        running = 0;
    }
}
exit(EXIT_SUCCESS);
}

```

## Shared memory

Shared memory allows two or more processes to access the same logical memory. Shared memory is efficient in transferring data between two running processes. Shared memory is a special range of addresses that is created by one process and the Shared memory appears in the address space of that process. Other processes then attach the same shared memory segment into their own address space. All processes can then access the memory location as if the memory had been allocated just like malloc. If one process writes to the shared memory, the changes immediately become visible to any other process that has access to the same shared memory.

The functions for shared memory are,

```
#include <sys/shm.h>
```

```
int shmget(key_t key, size_t size, int shmflg);  
void *shmat(int shm_id, const void *shm_addr, int shmflg);  
int shmctl(int shm_id, int cmd, struct shmid_ds *buf);  
int shmdt(const void *shm_addr);
```

The include files `sys/types.h` and `sys/ipc.h` are normally also required before `shm.h` is included.

### **shmget**

We create shared memory using the `shmget` function:

```
int shmget(key_t key, size_t size, int shmflg);
```

The argument `key` names the shared memory segment, and the `shmget` function returns a shared memory identifier that is used in subsequently shared memory functions. There's a special key-value, `IPC_PRIVATE`, that creates shared memory private to the process. The second parameter, `size`, specifies the amount of memory required in bytes. The third parameter, `shmflg`, consists of nine permission flags that are used in the same way as the mode flags for creating files. A special bit defined by `IPC_CREAT` must be bitwise ORed with the permissions to create a new shared memory segment. It's not an error to have the `IPC_CREAT` flag set and pass the key of an existing shared memory segment. The `IPC_CREAT` flag is silently ignored if it is not required.

The permission flags are very useful with shared memory because they allow a process to create shared memory that can be written by processes owned by the creator of the shared memory but only read by processes that other users have created. We can use this to provide efficient read-only access to data by placing it in shared memory without the risk of its being changed by other users.

If the shared memory is successfully created, `shmget` returns a non-negative integer, the shared memory identifier. On failure, it returns `-1`.

### **shmat**

When we first create a shared memory segment, it's not accessible by any process. To enable access to the shared memory, we must attach it to the address space of a process. We do this with the `shmat` function:

```
void *shmat(int shm_id, const void *shm_addr, int shmflg);
```

The first parameter, `shm_id`, is the shared memory identifier returned from `shmget`. The second parameter, `shm_addr`, is the address at which the shared memory is to be attached to the current process. This should almost always be a null pointer, which allows the system to choose the address at which the memory appears. The third parameter, `shmflg`,



is a set of bitwise flags. The two possible values are SHM\_RND, which, in conjunction with shm\_addr, controls the address at which the shared memory is attached, and SHM\_RDONLY, which makes the attached memory read-only. It's very rare to need to control the address at which shared memory is attached; you should normally allow the system to choose an address for you, as doing otherwise will make the application highly hardware-dependent. If the shmat call is successful, it returns a pointer to the first byte of shared memory. On failure -1 is returned.

The shared memory will have read or write access depending on the owner (the creator of the shared memory), the permissions, and the owner of the current process. Permissions on shared memory are similar to the permissions on files. An exception to this rule arises if shmflg & SHM\_RDONLY is true. Then the shared memory won't be writable, even if permissions would have allowed write access.

### **shmdt**

The shmdt function detaches the shared memory from the current process. It takes a pointer to the address returned by shmat. On success, it returns 0, on error -1. Note that detaching the shared memory doesn't delete it; it just makes that memory unavailable to the current process.

### **shmctl**

**int shmctl(int shm\_id, int command, struct shmid\_ds \*buf);**

The first parameter, shm\_id, is the identifier returned from shmget. The second parameter, command, is the action to take. It can take three values:

**Table 7.2 Values of shm\_id**

<b>Command</b>	<b>Description</b>
IPC_STAT	Sets the data in the shmid_ds structure to reflect the values associated with the shared memory.
IPC_SET	Sets the values associated with the shared memory to those provided in the shmid_ds data structure, if the process has permission to do so.
IPC_RMID	Deletes the shared memory segment.

The shmid\_ds structure has the following members:

```
struct shmid_ds {  
    uid_t shm_perm.uid;  
    uid_t shm_perm.gid;
```

```

    mode_t shm_perm.mode;
}

```

The third parameter, `buf`, is a pointer to the structure containing the modes and permissions for the shared memory. On success, it returns 0, on failure returns -1.

We will write a pair of programs `shm1.c` and `shm2.c`. The first will create a shared memory segment and display any data that is written into it. The second will attach into an existing shared memory segment and enters data into the shared memory segment.

First, we create a common header file to describe the shared memory we wish to pass around. We call this `shm_com.h`.

```

#define TEXT_SZ 2048

struct shared_use_st {
    int written_by_you;
    char some_text[TEXT_SZ];
};

```

//shm1.c – Consumer process

```

#include <unistd.h>
#include <stdlib.h>
#include <stdio.h>
#include <string.h>
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/shm.h>
#include "shm_com.h"

int main()
{
    int running = 1;
    void *shared_memory = (void *)0;
    struct shared_use_st *shared_stuff;
    int shmid;
    srand((unsigned int) getpid());
    shmid = shmget((key_t)1234, sizeof(struct shared_use_st), 0666 | IPC_CREAT);
    if (shmid == -1) {
        fprintf(stderr, "shmget failed\n");
        exit(EXIT_FAILURE);
    }
    shared_memory = shmat(shmid, (void *)0, 0);

```

```

if (shared_memory == (void *)-1) {
    fprintf(stderr, "shmat failed\n");
    exit(EXIT_FAILURE);
}
printf("Memory attached at %X\n", (int)shared_memory);
shared_stuff = (struct shared_use_st *)shared_memory;
shared_stuff->written_by_you = 0;
while(running) {
    if (shared_stuff->written_by_you) {
        printf("You wrote: %s", shared_stuff->some_text);
        sleep( rand() % 4 ); /* make the other process wait for us ! */
        shared_stuff->written_by_you = 0;
        if (strncmp(shared_stuff->some_text, "end", 3) == 0) {
            running = 0;
        }
    }
}
if (shmdt(shared_memory) == -1) {
    fprintf(stderr, "shmdt failed\n");
    exit(EXIT_FAILURE);
}
if (shmctl(shmid, IPC_RMID, 0) == -1) {
    fprintf(stderr, "shmctl(IPC_RMID) failed\n");
    exit(EXIT_FAILURE);
}
exit(EXIT_SUCCESS);
}

```

```

//shm2.c
#include <unistd.h>
#include <stdlib.h>
#include <stdio.h>
#include <string.h>
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/shm.h>
#include "shm_com.h"
int main()
{
    int running = 1;
    void *shared_memory = (void *)0;

```

```

struct shared_use_st *shared_stuff;
char buffer[BUFSIZ];
int shmid;
shmid = shmget((key_t)1234, sizeof(struct shared_use_st), 0666 | IPC_CREAT);
if (shmid == -1) {
    fprintf(stderr, "shmget failed\n");
    exit(EXIT_FAILURE);
}
shared_memory = shmat(shmid, (void *)0, 0);
if (shared_memory == (void *)-1) {
    fprintf(stderr, "shmat failed\n");
    exit(EXIT_FAILURE);
}
printf("Memory attached at %X\n", (int)shared_memory);
shared_stuff = (struct shared_use_st *)shared_memory;
while(running) {
    while(shared_stuff->written_by_you == 1) {
        sleep(1);
        printf("waiting for client...\n");
    }
    printf("Enter some text:");
    fgets(buffer, BUFSIZ, stdin);
    strncpy(shared_stuff->some_text, buffer, TEXT_SZ);
    shared_stuff->written_by_you = 1;
    if (strcmp(buffer, "end", 3) == 0) {
        running = 0;
    }
}
if (shmdt(shared_memory) == -1) {
    fprintf(stderr, "shmdt failed\n");
    exit(EXIT_FAILURE);
}
exit(EXIT_SUCCESS);
}

```

### Lab Exercises:

1. Process A wants to send a number to Process B. Once received, Process B has to check whether the number is palindrome or not. Write a C program to implement this interprocess communication using a message queue.
2. Implement a parent process, which sends an English alphabet to a child process using shared memory. The child process responds with the next English alphabet to the parent. The parent displays the reply from the Child.
3. Write two programs named Interface and Display for the following problem.

#### *Interface program*

This program, when run, gives a prompt to the user as “Enter your message:” When the user enters his/her message string and presses Enter, the program writes the message into the shared memory, tells the Display to start processing, and then prompts the user again for another message.

#### *Display program*

The process waits until a new message becomes available on the shared memory. Then it reads the contents of the memory and prints it on the screen. It also clears the contents of the shared memory when it has read the message.

4. Write a two-player 3x3 tic-tac-toe console game using shared memory.

### **Additional Exercises:**

1. Write a producer-consumer program in C in which producer writes a set of words into shared memory and then consumer reads the set of words from the shared memory. The shared memory needs to be detached and deleted after use.
2. Write a program which creates a message queue and writes message into queue which contains number of users working on the machine along with observed time in hours and minutes. This is repeated for every 10 minutes. Write another program which reads this information from the queue and calculates on average in each hour how many users are working.

### IPC – 3: DEADLOCK, LOCKING, SYNCHRONIZATION

In this lab, students will be able to:

- Synchronize various processes with the use of semaphore
- Understand how communication between two processes can take place with the help of a named pipe

For a multithreaded application spanning a single process or multiple processes to do useful work, some kind of common state must be shared between the threads. The degree of sharing that is necessary depends on the task. At one extreme, the only sharing necessary may be a single number that indicates the task to be performed. For example, a thread in a web server might be told only the port number to respond to. At the other extreme, a pool of threads might be passing information constantly among themselves to indicate what tasks are complete and what work is still to be completed.

#### Data Races

A data race occurs when multiple threads spanning a single process or multiple processes use the same data item and one or more of those threads are updating it.

Suppose there is a function update, which takes an integer pointer and updates the value of the content pointer by 4. If multiple threads call the function, then there is a possibility of a data race. If the current value of \*a is 10, then when two threads simultaneously call update function, the final value of \*a might be 14, instead of 18. To visualize this, we need to write the corresponding assembly language code for this function.

```
void update(int * a)
{
    *a = *a + 4;
}
```

Another situation might be when one thread is running, but the other thread has been context switched off of the processor. Imagine that the first thread has loaded the value of the variable **a** and then gets context switched off the processor. When it eventually runs again, the value of the variable a will have changed, and the final store of the restored thread will cause the value of the variable a to regress to an old value. The following code has a data race.

```
//race.c
#include <pthread.h>
int counter = 0;
void * func(void * params)
{
    counter++;
}
void main()
{
    pthread_t thread1, thread2;
```

```

    pthread_create(&thread1, 0, func, 0);
    pthread_create(&thread2, 0, func, 0);
    pthread_join(thread1, 0 );
    pthread_join(thread2, 0 );
}

```

### Using tools to detect data races

We can compile the above code using gcc, and then use Helgrind tool which is part of Valgrind suite to identify the data race.

```
$ gcc -g race.c -lpthread
```

```
$ valgrind --tool=helgrind ./a.out
```

### Avoiding Data Races

Although it is hard to identify data races, avoiding them can be very simple. The easiest way to do this is to place a synchronization lock around all accesses to that variable and ensure that before referencing the variable, the thread must acquire the lock.

### Synchronization Primitives:

#### Mutex Locks:

A mutex lock is a mechanism that can be acquired by only one thread at a time. Other threads that attempt to acquire the same mutex must wait until it is released by the thread that currently has it.

Mutex locks need to be initialized to the appropriate state by a call to `pthread_mutex_init()` or for statically defined mutexes by assignment with the `PTHREAD_MUTEX_INITIALIZER`. The call to `pthread_mutex_init()` takes an optional parameter that points to attributes describing the type of mutex required. Initialization through static assignment uses default parameters, as does passing in a null pointer in the call to `pthread_mutex_init()`.

Once a mutex is no longer needed, the resources it consumes can be freed with a call to `pthread_mutex_destroy()`.

```
#include <pthread.h>
```

```
...
```

```
pthread_mutex_t m1 = PTHREAD_MUTEX_INITIALIZER;
```

```
pthread_mutex_t m2;
```

```
pthread_mutex_init( &m2, 0 );
```

```
...
```

```
pthread_mutex_destroy( &m1 );
```

```
pthread_mutex_destroy( &m2 );
```

A thread can lock a mutex by calling `pthread_mutex_lock()`. Once it has finished with the mutex, the thread calls `pthread_mutex_unlock()`. If a thread calls `pthread_mutex_lock()` while another thread holds the mutex, the calling thread will wait, or *block*, until the other thread releases the mutex, allowing the calling thread to attempt to acquire the released mutex.

```
#include <pthread.h>
#include <stdio.h>
pthread_mutex_t mutex;
volatile int counter = 0;
void * count( void * param)
{
    for ( int i=0; i<100; i++)
    {
        pthread_mutex_lock(&mutex);
        counter++;
        printf("Count = %i\n", counter);
        pthread_mutex_unlock(&mutex);
    }
}
int main()
{
    pthread_t thread1, thread2;
    pthread_mutex_init( &mutex, 0 );
    pthread_create( &thread1, 0, count, 0 );
    pthread_create( &thread2, 0, count, 0 );
    pthread_join( thread1, 0 );
    pthread_join( thread2, 0 );
    pthread_mutex_destroy( &mutex );
    return 0;
}
```

### **Semaphores:**

A semaphore is a counting and signaling mechanism. One use for it is to allow threads access to a specified number of items. If there is a single item, then a semaphore is essentially the same as a mutex, but it is more commonly used in a situation where there are multiple items to be managed.

A semaphore is initialized with a call to `sem_init()`. This function takes three parameters. The first parameter is a pointer to the semaphore. The next is an integer to indicate whether the semaphore is shared between multiple processes or private to a single process. The final parameter is the value with which to initialize the semaphore. A semaphore created by a call to `sem_init()` is destroyed with a call to `sem_destroy()`.

The code below initializes a semaphore with a count of 10. The middle parameter of the call to `sem_init()` is zero, and this makes the semaphore private to the process; passing the value one rather than zero would enable the semaphore to be shared between multiple processes.

```
#include <semaphore.h>
int main()
{
```



```

    sem_t semaphore;
    sem_init( &semaphore, 0, 10 );
    ...
    sem_destroy( &semaphore );
}

```

The semaphore is used through a combination of two methods. The function `sem_wait()` will attempt to decrement the semaphore. If the semaphore is already zero, the calling thread will wait until the semaphore becomes nonzero and then return, having decremented the semaphore. The call to `sem_post()` will increment the semaphore. One more call, `sem_getvalue()`, will write the current value of the semaphore into an integer variable.

In the following program, an order is maintained in displaying Thread 1 and Thread 2. Try removing the semaphore and observe the output.

```

#include <pthread.h>
#include <stdio.h>
#include <semaphore.h>

sem_t semaphore;

void *func1( void * param )
{
    printf( "Thread 1\n" );
    sem_post( &semaphore );
}

void *func2( void * param )
{
    sem_wait( &semaphore );
    printf( "Thread 2\n" );
}

int main()
{
    pthread_t threads[2];
    sem_init( &semaphore, 0, 1 );
    pthread_create( &threads[0], 0, func1, 0 );
    pthread_create( &threads[1], 0, func2, 0 );
    pthread_join( threads[0], 0 );
    pthread_join( threads[1], 0 );
    sem_destroy( &semaphore );
}

```

### **Solution to Producer-Consumer problem**

```

#include<stdio.h>
#include<pthread.h>
#include<semaphore.h>

```

```

int buf[5],f,r;
sem_t mutex,full,empty;
void *produce(void *arg)
{
    int i;
    for(i=0;i<10;i++)
    {
        sem_wait(&empty);
        sem_wait(&mutex);
        printf("produced item is %d\n",i);
        buf[(++r)%5]=i;
        sleep(1);
        sem_post(&mutex);
        sem_post(&full);
        printf("full %u\n",full);
    }
}
void *consume(void *arg)
{
    int item,i;
    for(i=0;i<10;i++)
    {
        sem_wait(&full);
        printf("full %u\n",full);
        sem_wait(&mutex);
        item=buf[(++f)%5];
        printf("consumed item is %d\n",item);
        sleep(1);
        sem_post(&mutex);
        sem_post(&empty);
    }
}
main()
{
    pthread_t tid1,tid2;
    sem_init(&mutex,0,1);
    sem_init(&full,0,1);
    sem_init(&empty,0,5);
    pthread_create(&tid1,NULL,produce,NULL);
    pthread_create(&tid2,NULL,consume,NULL);
    pthread_join(tid1,NULL);
    pthread_join(tid2,NULL);
}

```

### **Solution to First Readers-Writers Problem using semaphores:**

- The reader processes share the following data structures:

```
semaphore mutex , wrt;  
int readcount;
```

- The binary semaphores mutex and wrt are initialized to 1; readcount is initialized to 0;
- Semaphore wrt is common to both reader and writer process
  - wrt functions as a mutual exclusion for the writers
  - It is also used by the first or last reader that enters or exits the critical section
  - It is not used by readers who enter or exit while other readers are in their critical section
- The readcount variable keeps track of how many processes are currently reading the object
- The mutex semaphore is used to ensure mutual exclusion when readcount is updated

The structure of a writer process

```
do {  
    wait(wrt);  
    . . .  
    // writing is performed  
    . . .  
    signal(wrt);  
} while (TRUE);
```


The structure of a reader process

```
do {  
    wait(mutex);  
    readcount++;  
    if (readcount == 1)  
        wait(wrt);  
    signal(mutex);  
    . . .  
    // reading is performed  
    . . .  
    wait(mutex);  
    readcount--;  
    if (readcount == 0)  
        signal(wrt);  
    signal(mutex);  
} while (TRUE);
```

### **The Dining Philosophers Problem:**

- Five philosophers sit at a round table - thinking and eating
- Each philosopher has one chopstick
  - five chopsticks total
- A philosopher needs two chopsticks to eat
  - philosophers must share chopsticks to eat
- No interaction occurs while thinking

The situation of the dining philosophers is shown in Fig. 7.1



**Figure 8.1 Dining Philosophers Problem**

**Lab Exercises:**

1. Modify the above Producer-Consumer program so that, a producer can produce at the most 10 items more than what the consumer has consumed.
2. Write a C program for the first readers-writers problem using semaphores.
3. Write a Code to access a shared resource which causes deadlock using improper use of semaphore.
4. Write a program using semaphore to demonstrate the working of sleeping barber problem.

**Additional Exercises:**

1. Write a C program for Dining-Philosophers problem using monitors.
2. Demonstrate the working of counting semaphore.

## PROGRAMS ON THREADS

### Objectives:

In this lab, the student will be able to:

- Understand the concepts of multithreading
- Grasp the execution of the different processes concerning multithreading

A process will start with a single thread which is called the main thread or master thread. Calling `pthread_create()` creates a new thread. It takes the following parameters.

- A pointer to a `pthread_t` structure. The call will return the handle to the thread in this structure.
- A pointer to a `pthread` attributes structure, which can be a null pointer if the default attributes are to be used. The details of this structure will be discussed later.
- The address of the routine to be executed.
- A value or pointer to be passed into the new thread as a parameter.

```
#include <pthread.h>
```

```
#include <stdio.h>
```

```
void* thread_code( void * param )
{
    printf( "In thread code\n" );
}
int main()
{
    pthread_t thread;
    pthread_create(&thread, 0, &thread_code, 0 );
    printf("In main thread\n" );
}
```

In this example, the main thread will create a second thread to execute the routine `thread_code()`, which will print one message while the main thread prints another. The call to create the thread has a value of zero for the attributes, which gives the thread default attributes. The call also passes the address of a `pthread_t` variable for the function to store a handle to the thread. The return value from the `thread_create()` call is zero if the call is successful; otherwise, it returns an error condition.

### Thread termination :

Child threads terminate when they complete the routine they were assigned to run. In the above example, child thread will terminate when it completes the routine `thread_code()`.

The value returned by the routine executed by the child thread can be made available to the main thread when the main thread calls the routine `pthread_join()`.

The `pthread_join()` call takes two parameters. The first parameter is the handle of the thread that is to be waited for. The second parameter is either zero or the address of a pointer to a void, which will hold the value returned by the child thread.

The resources consumed by the thread will be recycled when the main thread calls `pthread_join()`. If the thread has not yet terminated, this call will wait until the thread terminates and then free the assigned resources.

```
#include <pthread.h>
#include <stdio.h>

void* thread_code( void * param )
{
    printf( "In thread code\n" );
}

int main()
{
    pthread_t thread;
    pthread_create( &thread, 0, &thread_code, 0 );
    printf( "In main thread\n" );
    pthread_join( thread, 0 );
}
```

Another way a thread can terminate is to call the routine `pthread_exit()`, which takes a single parameter—either zero or a pointer—to void. This routine does not return and instead terminates the thread. The parameter passed into the `pthread_exit()` call is returned to the main thread through the `pthread_join()`. The child threads do not need to explicitly call `pthread_exit()` because it is implicitly called when the thread exits.

### Passing Data to and from Child Threads

In many cases, it is important to pass data into the child thread and have the child thread return status information when it completes. To pass data into a child thread, it should be cast as a pointer to void and then passed as a parameter to `pthread_create()`.

```
for ( int i=0; i<10; i++ )
    pthread_create( &thread, 0, &thread_code, (void *)i );
```

Following is a program where the main thread passes a value to the Pthread and the thread returns a value to the main thread.

```
#include <pthread.h>
#include <stdio.h>

void* child_thread( void * param )
{
    int id = (int)param;
    printf( "Start thread %i\n", id );
    return (void *)id;
}
```

```

int main()
{
    pthread_t thread[10];
    int return_value[10];
    for ( int i=0; i<10; i++ )
    {
        pthread_create( &thread[i], 0, &child_thread, (void*)i );
    }
    for ( int i=0; i<10; i++ )
    {
        pthread_join( thread[i], (void**)&return_value[i] );
        printf( "End thread %i\n", return_value[i] );
    }
}

```

### Setting the Attributes for Pthreads

The attributes for a thread are set when the thread is created. To set the initial thread attributes, first, create a thread attributes structure, and then set the appropriate attributes in that structure, before passing the structure into the `pthread_create()` call.

```

#include <pthread.h>
...
int main()
{
    pthread_t thread;
    pthread_attr_t attributes;
    pthread_attr_init( &attributes );
    pthread_create( &thread, &attributes, child_routine, 0 );
}

```

### Lab Exercises:

1. Write a multithreaded program that generates the Fibonacci series. The program should work as follows: The user will enter on the command line the number of Fibonacci numbers that the program is to generate. The program then will create a separate thread that will generate the Fibonacci numbers, placing the sequence in data that is shared by the threads (an array is probably the most convenient data structure). When the thread finishes execution the parent will output the sequence generated by the child thread. Because the parent thread cannot begin outputting the Fibonacci sequence until the child thread finishes, this will require having the parent thread wait for the child thread to finish.
2. Write a multithreaded program that calculates the summation of non-negative integers in a separate thread and passes the result to the main thread.
3. Write a multithreaded program for generating prime numbers from a given starting number to the given ending number.
4. Write a multithreaded program that performs the sum of even numbers and odd numbers in an input array. Create a separate thread to perform the sum of even numbers and odd numbers. The parent thread has to wait until both the threads are done.

### Additional Exercises:

1. Write a multithreaded program for matrix multiplication.
2. Write a multithreaded program for finding row sum and column sum

## **MEMORY AND DATA MANAGEMENT**

### **Objectives:**

In this lab, the student will be able to:

- Understand how to use dynamic memory allocation
- Learn the working of Demand Paged Virtual Memory

### **Simple Memory Allocation**

You allocate memory using the malloc call in the standard C library:

```
#include <stdlib.h>
```

```
void *malloc(size_t size);
```

Notice that Linux (following the X/Open specification) differs from some UNIX implementations by not requiring a special malloc.h include file. Note also that the size parameter that specifies the number of bytes to allocate isn't a simple int, although it's usually an unsigned integer type.

You can allocate a great deal of memory on most Linux systems. Let's start with a very simple program, but one that would defeat old MS-DOS-based programs, because they cannot access memory outside the base 640K memory map of PCs.

### **Sample Program (memory1.c) :**

```
#include <unistd.h>
```

```
#include <stdlib.h>
```

```
#include <stdio.h>
```

```
#define A_MEGABYTE (1024 * 1024)
```

```
int main()
```

```
{
```

```
char *some_memory;
```

```
int megabyte = A_MEGABYTE;
```

```
int exit_code = EXIT_FAILURE;
```

```
some_memory = (char *)malloc(megabyte);
```

```
if (some_memory != NULL) {
```

```
    sprintf(some_memory, "Hello World\n");
```

```
    printf("%s", some_memory);
```

```
    exit_code = EXIT_SUCCESS;
```

```
}
```

```
exit(exit_code);
```

```
}
```

When you run this program, it gives the following output:

```
$ ./memory1
```

```
Hello World
```

### **How It Works**

This program asks the malloc library to give it a pointer to a megabyte of memory. You check to ensure that malloc was successful and then use some of the memory to show that it exists. When you run the program, you should see Hello World printed out, showing that malloc did indeed return the megabyte of usable memory. We don't check that all of the megabytes are present; we have to put some trust in the malloc code!

Notice that because malloc returns a void \* pointer, you cast the result to the char \* that you need.

The malloc function is guaranteed to return memory that is aligned so that it can be cast to a pointer of any type. The simple reason is that most current Linux systems use 32-bit integers and 32-bit pointers for pointing to memory, which allows you to specify up to 4 gigabytes. This



ability to address directly with a 32-bit pointer, without needing segment registers or other tricks, is termed a flat 32-bit memory model.

### **Allocating Lots of Memory**

Now that you've seen Linux exceed the limitations of the MS-DOS memory model, let's give it a more difficult problem. The next program asks to allocate somewhat more memory than is physically present in the machine, so you might expect malloc to start failing somewhere a little short of the actual amount of memory present, because the kernel and all the other running processes are using some memory.

### **Sample Program (memory2.c) :**

Asking for All Physical Memory

With memory2.c , we're going to ask for more than the machine's physical memory. You should adjust the define PHY\_MEM\_MEGS depending on your physical machine:

```
#include <unistd.h>
#include <stdlib.h>
#include <stdio.h>
#define A_MEGABYTE (1024 * 1024)
#define PHY_MEM_MEGS
1024 /* Adjust this number as required */
int main()
{
    char *some_memory;
    size_t size_to_allocate = A_MEGABYTE;
    int megs_obtained = 0;
    while (megas_obtained < (PHY_MEM_MEGS * 2)) {
        some_memory = (char *)malloc(size_to_allocate);
        if (some_memory != NULL) {
            megs_obtained++;
            sprintf(some_memory, "Hello World");
            printf("%s - now allocated %d Megabytes\n", some_memory, megs_obtained);
        }
        else {
            exit(EXIT_FAILURE);
        }
    }
    exit(EXIT_SUCCESS);
}
```

The output, somewhat abbreviated, is as follows:

```
$ ./memory2
Hello World
Hello World
...
Hello World
Hello World
- now allocated 1 Megabytes
- now allocated 2 Megabytes
- now allocated 2047 Megabytes
- now allocated 2048 Megabytes
```

## How It Works

The program is very similar to the previous example. It simply loops, asking for more and more memory, until it has allocated twice the amount of memory you said your machine had when you adjusted the define `PHY_MEM_MEGS`. The surprise is that it works at all because we appear to have created a program that uses every single byte of physical memory on the author's machine. Notice that we use the `size_t` type for our call to `malloc`.

The other interesting feature is that, at least on this machine, it ran the program in the blink of an eye. So not only have we used up all the memory, but we've done it very quickly indeed.

Let's investigate further and see just how much memory we can allocate on this machine with `memory3.c`. Since it's now clear that Linux can do some very clever things with memory requests, we'll allocate memory just 1K at a time and write to each block that we obtain.

## Demand Paged Virtual Memory

### Sample Program (`memory3.c`) :

Available Memory

```
#include <unistd.h>
#include <stdlib.h>
#include <stdio.h>
#define ONE_K (1024)
int main()
{
    char *some_memory;
    int size_to_allocate = ONE_K;
    int megs_obtained = 0;
    int ks_obtained = 0;
    while (1) {
        for (ks_obtained = 0; ks_obtained < 1024; ks_obtained++) {
            some_memory = (char *)malloc(size_to_allocate);
            if (some_memory == NULL) exit(EXIT_FAILURE);
            sprintf(some_memory, "Hello World");
        }
        megs_obtained++;
        printf("Now allocated %d Megabytes\n", megs_obtained);
    }
    exit(EXIT_SUCCESS);
}
```

### Output:

```
$ ./memory3
Now allocated 1 Megabytes
...
Now allocated 1535 Megabytes
Now allocated 1536 Megabytes
Out of Memory: Killed process 2365
Killed
```

and then the program ends. It also takes quite a few seconds to run, and slows down significantly around the same number as the physical memory in the machine, and exercises the hard disk quite noticeably. However, the program has allocated, and accessed, more

memory than this author physically has in his machine at the time of writing. Finally, the system protects itself from this rather aggressive program and kills it. On some systems, it may simply exit quietly when malloc fails.

### How It Works

The application's allocated memory is managed by the Linux kernel. Each time the program asks for memory or tries to read or write to memory that it has allocated, the Linux kernel takes charge and decides how to handle the request. Initially, the kernel was simply able to use free physical memory to satisfy the application's request for memory, but once physical memory was full, it started using what's called swap space. On Linux, this is a separate disk area allocated when the system was installed. If you're familiar with Windows, the Linux swap space acts a little like the hidden Windows swap file. However, unlike Windows, there are no local heap, global heap, or discardable memory segments to worry about in code — the Linux kernel does all the management for you. The kernel moves data and program code between physical memory and the swap space so that each time you read or write memory, the data always appears to have been in physical memory, wherever it was located before you attempted to access it.

- In more technical terms, Linux implements a **demand paged virtual memory** system. All memory seen by user programs is virtual; that is, it doesn't exist at the physical address the program uses. Linux divides all memory into pages, commonly 4,096 bytes per page. When a program tries to access memory, a virtual-to-physical translation is made, although how this is implemented and the time it takes depend on the particular hardware you're using. When the access is to memory that isn't physically resident, a page fault results and control is passed to the kernel. The Linux kernel checks the address being accessed and, if it's a legal address for that program, determines which page of physical memory to make available. It then either allocates it, if it has never been written before or if it has been stored on the disk in the swap space, reads the memory page containing the data into physical memory (possibly moving an existing page out to disk). Then, after mapping the virtual memory address to match the physical address, it allows the user program to continue. Linux applications don't need to worry about this activity because the implementation is all hidden in the kernel. Eventually, when the application exhausts both the physical memory and the swap space, or when the maximum stack size is exceeded, the kernel finally refuses the request for further memory and may pre-emptively terminate the program.

This “killing the process” behavior is different from early versions of Linux and many other flavors of UNIX, where malloc simply fails. It's termed the “out of memory (OOM) killer,” and although it may seem rather drastic, it is a good compromise between letting processes allocate memory rapidly and efficiently and having the Linux kernel protect itself from a total lack of resources, which is a serious issue.

So what does this mean to the application programmer? It's all good news. Linux is very good at managing memory and will allow applications to use very large amounts of memory and even very large single blocks of memory. However, you must remember that allocating two blocks of memory won't result in a single continuously addressable block of memory. What you get is what you ask for: two separate blocks of memory.

Does this limitless supply of memory, followed by the preemptive killing of the process, mean that there's no point in checking the return from malloc? No. One of the most common problems in C programs using dynamically allocated memory is writing beyond the end of an allocated block. When this happens, the program may not terminate immediately, but you have probably overwritten some data used internally by the malloc library routines.

Usually, the result is that future calls to malloc may fail, not because there's no memory to allocate, but because the memory structures have been corrupted. These problems can be quite difficult to track down, and in programs the sooner the error is detected, the better the chances of tracking down the cause.

### **Abusing Memory**

Suppose you try to do “bad” things with memory. In this exercise, you allocate some memory and then attempt to write past the end, in memory4.c .

#### **Sample Program (memory4.c) :**

```
#include <stdlib.h>
#define ONE_K (1024)
int main()
{
    char *some_memory;
    char *scan_ptr;
    some_memory = (char *)malloc(ONE_K);
    if (some_memory == NULL) exit(EXIT_FAILURE);
    scan_ptr = some_memory;
    while(1) {
        *scan_ptr = '\0';
        scan_ptr++;
    }
    exit(EXIT_SUCCESS);
}
```

#### **The output is simply**

```
$ /memory4
Segmentation fault
```

### **How It Works**

The Linux memory management system has protected the rest of the system from this abuse of memory. To ensure that one badly behaved program (this one) can't damage any other programs, Linux has terminated it. Each running program on a Linux system sees its memory map, which is different from every other program's. Only the operating system knows how physical memory is arranged, and not only manages it for user programs but also protects user programs from each other.

### **The Null Pointer**

Unlike MS-DOS, but more like newer flavors of Windows, modern Linux systems are very protective about writing or reading from the address referred to by a null pointer, although the actual behavior is implementation-specific.

#### **Sample Program (memory5a.c) :**

```
#include <unistd.h>
#include <stdlib.h>
#include <stdio.h>
int main()
{
    char *some_memory = (char *)0;
    printf("A read from null %s\n", some_memory);
    sprintf(some_memory, "A write to null\n");
    exit(EXIT_SUCCESS);
}
```

### **The output is**

```
$ ./memory5a
A read from null (null)
Segmentation fault
```

### **How It Works**

The first `printf` attempts to print out a string obtained from a null pointer; then the `sprintf` attempts to write to a null pointer. In this case, Linux (in the guise of the GNU “C” library) has been forgiving about the read and has simply provided a “magic” string containing the characters `( n u l l ) \0`. It hasn’t been so forgiving about the write and has terminated the program. This can sometimes help track down program bugs.

If you try this again but this time don’t use the GNU “C” library, you’ll discover that reading from location zero is not permitted. Here is `memory5b.c` :

### **Sample Program (memory5b.c) :**

```
#include <unistd.h>
#include <stdlib.h>
#include <stdio.h>
int main()
{
    char z = *(const char *)0;
    printf(“I read from location zero\n”);
    exit(EXIT_SUCCESS);
}
```

The output is  
`$ ./memory5b`  
Segmentation fault

This time you attempt to read directly from location zero. There is no GNU libc library between you and the kernel now, and the program is terminated. Note that some versions of UNIX do permit reading from location zero, but Linux doesn’t.

### **Freeing Memory**

Up to now, we’ve been simply allocating memory and then hoping that when the program ends, the memory we’ve used hasn’t been lost. Fortunately, the Linux memory management system is quite capable of reliably ensuring that memory is returned to the system when a program ends. However, most programs don’t simply want to allocate some memory, use it for a short period, and then exit. A much more common use is dynamically using memory as required.

Programs that use memory on a dynamic basis should always release unused memory back to the malloc memory manager using the `free` call. This enables separate blocks to be remerged and enables the malloc library to look after memory, rather than have the application manage it. If a running program (process) uses and then frees memory, that free memory remains allocated to the process. Behind the scenes, Linux is managing the blocks of memory the programmer is using as a set of physical “pages,” usually 4K bytes each, in memory. However, if a page of memory is not being used, then the Linux memory manager will be able to move it from physical memory to swap space (termed paging), where it has little impact on the use of resources. If the program tries to access data inside the memory page that has been moved to swap space, then Linux will very briefly suspend the program, move the memory page back from swap space into physical memory again, and then allow the program to continue, just as though the data had been in memory all along.

```
#include <stdlib.h>
void free(void *ptr_to memory);
```

A call to free should be made only with a pointer to memory allocated by a call to malloc, calloc, or realloc.

#### Sample Program (memory6.c) :

```
#include <stdlib.h>
#include <stdio.h>
#define ONE_K (1024)
int main()
{
    char *some_memory;
    int exit_code = EXIT_FAILURE;
    some_memory = (char *)malloc(ONE_K);
    if (some_memory != NULL) {
        free(some_memory);
        printf("Memory allocated and freed again\n");
        exit_code = EXIT_SUCCESS;
    }
    exit(exit_code);
}
```

The output is

```
$ ./memory6
```

```
Memory allocated and freed again
```

#### How It Works

This program simply shows how to call free with a pointer to some previously allocated memory.

Remember that once you've called free on a block of memory, it no longer belongs to the process. It's not being managed by the malloc library. **Never try to read or write memory after calling free on it.**

#### Other Memory Allocation Functions

Two other memory allocation functions are not used as often as malloc and free : calloc and realloc .

The prototypes are

```
#include <stdlib.h>
void *calloc(size_t number_of_elements, size_t element_size);
void *realloc(void *existing_memory, size_t new_size);
```

Although calloc allocates memory that can be freed with free , it has somewhat different parameters from malloc : It allocates memory for an array of structures and requires the number of elements and the size of each element as its parameters. The allocated memory is filled with zeros; and if calloc is successful, a pointer to the first element is returned. Like malloc , subsequent calls are not guaranteed to return contiguous space, so you can't enlarge an array created by calloc by simply calling calloc again and expecting the second call to return memory appended to that returned by the first call.

The realloc function changes the size of a block of memory that has been previously allocated. It's passed a pointer to some memory previously allocated by malloc , calloc , or realloc and

resizes it up or down as requested. The realloc function may have to move data around to achieve this, so it's important to ensure that once the memory has been reallocated, you always use the new pointer and never try to access the memory using pointers set up before realloc was called.

Another problem to watch out for is that realloc returns a null pointer if it has been unable to resize the memory. This means that in some applications you should avoid writing code like this:

```
my_ptr = malloc(BLOCK_SIZE);
....
my_ptr = realloc(my_ptr, BLOCK_SIZE * 10);
```

If realloc fails, then it returns a null pointer; my\_ptr will point to null, and the original memory allocated with malloc can no longer be accessed via my\_ptr . It may, therefore, be to your advantage to request the new memory first with malloc and then copy data from the old block to the new block using memcpy before freeing the old block. On error, this would allow the application to retain access to the data stored in the original block of memory, perhaps while arranging a clean termination of the program.

### Lab Exercises

1. If you wish to implement Best Fit, First Fit, Next Fit, or Worst Fit memory allocation policy, it is probably best to do this by describing the memory as a structure in a linked list:

```
struct mab {
    int offset;
    int size;
    int allocated;
    struct mab * next;
    struct mab * prev;
};
typedef struct mab Mab;
typedef Mab * MabPtr;
```

Either way, the following set of prototypes give a guide as to the functionality you will need to provide:

**MabPtr memChk(MabPtr m, int size); // check if memory available**

**MabPtr memAlloc(MabPtr m, int size); // allocate a memory block**

**MabPtr memFree(MabPtr m); // free memory block**


**MabPtr memMerge(MabPtr m); // merge two memory blocks**

**MabPtr memSplit(MabPtr m, int size); // split a memory block**

2. Write a C program using Malloc for implementing Multilevel feedback queue using three queues with each of them working with different scheduling policies

3. We have five segments numbered 0 through 4. The segments are stored in physical memory as shown in the following Fig 10.3. Write a C program to create segment table. Write methods for converting logical address to physical address. Compute the physical address for the following.

(i) 53 byte of segment 2 (ii) 852 byte of segment 3 (iii) 1222 byte of segment 0



**Figure 9.1: Physical memory**

4. Write a C program to simulate LRU approximation page replacement using second chance algorithm. Find the total number of page faults and hit ratio for the algorithm.

#### **Additional Exercises:**

1. Implement the same concept for the Buddy system.
2. How do you resize and release memory using realloc ?



## DEADLOCK AND DISK MANAGEMENT

### Objectives:

In this lab, students will be able to

- Understand the problem of deadlock and ways to manage it
- Understand find the details of underlying operating systems disk space and file system information

### DEADLOCK MANAGEMENT

#### The deadlock problem:

A set of blocked processes each holding a resource and waiting to acquire a resource held by another process in the set.




Figure 11.1: Deadlock Situation

Above Fig. 11.1 shows, a situation of deadlock where Process P1 Waiting for resource R2, which is held with process P2 and in the meantime, Process P2 is waiting for resource R1, which is held with process P1. Neither P1 nor P2 can proceed their execution until their needed resources are fulfilled forming a cyclic wait. It is the deadlock situation among processes as both are not progressed. In a single instance of resource type, a cyclic wait is always a deadlock. Consider Figure 11.2 below, the situation with 4 processes P1, P2, P3 and P4 and 2 resources R1 and R2 both are of two instances. Here, there is no deadlock even though the cycle exists between processes P1 and P3. Once P2 finishes its job, 1 instance of resource will be available which can be accessed by process P1, which turns request edge to assignment edge, thereby removing cyclic-wait. So, in multiple instances of resource type, the cyclic-wait need not be deadlock.




Figure 11.2: Cyclic-wait but no deadlock

## Methods for Handling Deadlocks:

### (i) Deadlock Avoidance:

The deadlock avoidance algorithm dynamically examines the resource-allocation state to ensure that there can never be a circular-wait condition. Resource-allocation *state* is defined by the number of available and allocated resources, and the maximum demands of the processes.

#### Safe State:

System is in safe state if there exists a safe sequence of all processes. **Sequence** of processes  $\langle P_1, P_2, \dots, P_n \rangle$  is **safe** if for each  $P_i$ , the resources that  $P_i$  can still request can be satisfied by currently available resources + resources held by all the  $P_j$ , with  $j < i$ .

- If  $P_i$  resource needs are not immediately available, then  $P_i$  can wait until all  $P_j$  have finished.
- When  $P_j$  is finished,  $P_i$  can obtain needed resources, execute, return allocated resources, and terminate.
- When  $P_i$  terminates,  $P_{i+1}$  can obtain its needed resources, and so on.

If a system is in safe state  $\Rightarrow$  no deadlocks.

If a system is in unsafe state  $\Rightarrow$  possibility of deadlock.

Avoidance  $\Rightarrow$  ensure that a system will never enter an unsafe state.

## DISK MANAGEMENT

### System calls related to disk

#### df command

The 'df' (Disk Free) command is an inbuilt utility to find the available and the disk usage space on Linux servers/storage.

The following Table 11.1 provides an overview of the options of df command in Linux.

**Table 11.1 Options of df command**

Options	Description
-a	-all: includes real files and virtual files like pseudo, pro, sysfs, lxc
-h	It prints the sizes in the human-readable format in power of 1024 (eg: 1K 1M 1G)
-H	--si : likewise '-h' but here in power of 1000
-i	--inodes: correspondence the inode details
-k	--block-size=1K display the disk space
-l	--local display local file systems only
-m	--megabytes display the disk space
-t	--type=TYPE To filter a particular file system type
-T	--print-type List the file system types
-x	--exclude-type=TYPE To exclude a particular file system type

#### 1. How to check the details of disk space used in each file system?

```
# df
```

Output:

Filesystem	1K-blocks	Used	Available	Use%	Mounted on
/dev/sda2	164962420	142892148	15301196	91%	/
udev	10240	0	10240	0%	/dev
tmpfs	3291620	329084	2962536	10%	/run
tmpfs	8229048	0	8229048	0%	/dev/shm
tmpfs	5120	0	5120	0%	/run/lock
tmpfs	8229048	0	8229048	0%	/sys/fs/cgroup
/dev/sda1	97167	76552	15598	84%	/boot
tmpfs	1645812	0	1645812	0%	/run/user/301703
tmpfs	1645812	0	1645812	0%	/run/user/301677
tmpfs	1645812	0	1645812	0%	/run/user/301483

Note: Using 'df' command without any option/parameter will display all the partitions and the usage information of the disk space. The result of the above command contains 6 columns which are explained here below:

Filesystem	-->	Mount Point Name
1K-blocks	-->	Available total space counted in 1kB (1000 bytes)
Used	-->	Used block size
Available	-->	Free blocks size
Use%	-->	Usage on percentage-wise
Mounted on	-->	Show the path of the mounted point

# df -k

Note: Even using the '-k' option also provides the same output as the default 'df' command. Both outputs provide the same data usage of file systems in block size which is measured in 1024 bytes.

## 2. How to sum up the total of the disk space usage?

```
# df -h --total
```

Output:

Filesystem	Size	Used	Avail	Use%	Mounted on
/dev/sda2	158G	137G	15G	91%	/
udev	10M	0	10M	0%	/dev
tmpfs	3.2G	322M	2.9G	11%	/run
tmpfs	7.9G	0	7.9G	0%	/dev/shm
tmpfs	5.0M	0	5.0M	0%	/run/lock
tmpfs	7.9G	0	7.9G	0%	/sys/fs/cgroup
/dev/sda1	95M	75M	16M	84%	/boot
tmpfs	1.6G	0	1.6G	0%	/run/user/301703
tmpfs	1.6G	0	1.6G	0%	/run/user/301483
tmpfs	1.6G	0	1.6G	0%	/run/user/301613
tmpfs	1.6G	0	1.6G	0%	/run/user/301677
total	183G	137G	40G	78%	-

Note: using '--total' along with '-h' will sum up the total disk usage of all the file systems.

### 3. How to list the Inodes information of all file systems?

```
# df -i
```

Output:

Filesystem	Inodes	IUsed	IFree	IUse%	Mounted on
/dev/sda2	10240000	2491518	7748482	25%	/
udev	2054985	305	2054680	1%	/dev
tmpfs	2057262	767	2056495	1%	/run
tmpfs	2057262	1	2057261	1%	/dev/shm
tmpfs	2057262	11	2057251	1%	/run/lock
tmpfs	2057262	13	2057249	1%	/sys/fs/cgroup
/dev/sda1	25168	336	24832	2%	/boot
tmpfs	2057262	4	2057258	1%	/run/user/301703
tmpfs	2057262	4	2057258	1%	/run/user/301483
tmpfs	2057262	4	2057258	1%	/run/user/301613
tmpfs	2057262	4	2057258	1%	/run/user/301677

Note: Using '-i' will list the information about the Inodes of all the filesystem.

Disk scheduling algorithms for the disk structure as shown in the Figure 10.1 are used by operating systems to determine the order in which read and write operations are performed on a disk. The main goal of these algorithms is to minimize the disk head movements and optimize the overall disk performance.


Here are some common disk scheduling algorithms:

1. First-Come, First-Served (FCFS): In this algorithm, the disk requests are executed in the order they arrive. The disk head moves from its current position to the requested track, serving the requests sequentially. FCFS is simple but can result in poor performance due to the phenomenon called the "elevator effect."
2. Shortest Seek Time First (SSTF): This algorithm selects the request that requires the least disk head movement from the current position. It minimizes the seek time by serving the closest request first. SSTF provides better performance compared to FCFS, but it may cause starvation for requests located farther from the current position.
3. SCAN: Also known as the elevator algorithm, SCAN moves the disk head in one direction, serving requests along the way until it reaches the end of the disk. Then, it changes direction and serves requests in the opposite direction. SCAN provides a fair servicing order for all requests but may cause delays for requests located at the extremes.
4. C-SCAN: Similar to SCAN, C-SCAN moves the disk head in one direction, serving requests until the end of the disk is reached. However, instead of changing direction, it immediately jumps to the opposite end and starts servicing requests from there. This algorithm avoids the delays caused by SCAN for requests located at the extremes.

5. **LOOK:** LOOK is an improvement over SCAN. Instead of moving to the end of the disk, LOOK changes direction as soon as there are no pending requests in the current direction. This reduces the head movement and improves the average seek time.

6. **C-LOOK:** C-LOOK combines the advantages of C-SCAN and LOOK. It moves the disk head in one direction, serving requests until the last request in that direction. Then, it jumps to the opposite end and starts servicing requests from there without reversing direction. C-LOOK reduces head movement and provides better performance than both SCAN and LOOK.

The First-Come, First-Served (FCFS) algorithm is one of the simplest disk scheduling algorithms. It operates on the principle that requests are serviced in the order they arrive. In the context of disk scheduling, FCFS refers to the order in which read and write requests are executed on the disk.



**Figure 11.3: Disk structure**

### Sample Program:

Simulating FCFS Disk Scheduling algorithm

```
#include<stdio.h>
#include<conio.h>
void main()
{
    int queue[100],n,head,i,j,k,seek=0,diff;
    float avg;
    // clrscr();
    printf("*** FCFS Disk Scheduling Algorithm ***\n");
    printf("Enter the size of Queue\t");
    scanf("%d",&n);
    printf("Enter the Queue\t");
```

```

for(i=1;i<=n;i++)
{
    scanf("%d",&queue[i]);
}
printf("Enter the initial head position\t");
scanf("%d",&head);
queue[0]=head;
printf("\n");
for(j=0;j<=n-1;j++)
{
    diff=abs(queue[j+1]-queue[j]);
    seek+=diff;
    printf("Move from %d to %d with Seek %d\n",queue[j],queue[j+1],diff);
}
printf("\nTotal Seek Time is %d\t",seek);
avg=seek/(float)n;
printf("\nAverage Seek Time is %f\t",avg);
getch();
}

```

### Output

```

*** FCFS Disk Scheduling Algorithm ***
Enter the size of Queue 8
Enter the Queue 98 183 37 122 14 124 65 67
Enter the initial head position 53

Move from 53 to 98 with Seek 45
Move from 98 to 183 with Seek 85
Move from 183 to 37 with Seek 146
Move from 37 to 122 with Seek 85
Move from 122 to 14 with Seek 108
Move from 14 to 124 with Seek 110
Move from 124 to 65 with Seek 59
Move from 65 to 67 with Seek 2

Total Seek Time is 640
Average Seek Time is 80.000000

```

### **Lab Exercises:**

1. Consider the following snapshot of the system. Write C program to implement Banker's algorithm for deadlock avoidance. The program has to accept all inputs from the user. Assume the total number of instances of A,B and C are 10,5 and 7 respectively.

	<u>Allocation</u>	<u>Max</u>	<u>Available</u>
	<u>A B C</u>	<u>A B C</u>	<u>A B C</u>
$P_0$	0 1 0	7 5 3	3 3 2
$P_1$	2 0 0	3 2 2	
$P_2$	3 0 2	9 0 2	
$P_3$	2 1 1	2 2 2	
$P_4$	0 0 2	4 3 3	

- (a) What is the content of the matrix *Need*?
- (b) Is the system in a safe state?

- (c) If a request from process P1 arrives for (1, 0, 2), can the request be granted immediately? Display the updated Allocation, Need and Available matrices.
- (d) If a request from process P4 arrives for (3, 3, 0), can the request be granted immediately?
- (e) If a request from process P0 arrives for (0, 2, 0), can the request be granted immediately?
- Write a multithreaded program that implements the banker's algorithm. Create  $n$  threads that request and release resources from the bank. The banker will grant the request only if it leaves the system in a safe state. You may write this program using **pthread**. It is important that shared data be safe from concurrent access. To ensure safe access to shared data, you can use mutex locks, which are available in the pthreads libraries.
  - Simulate implementation of Disk Scheduling Algorithms: FCFS, SSTF using a structure DSA. An DSA contains the request ID, arrival timestamp, cylinder, address, and the ID of the process that posted the request.
 

```
struct DSA {
    int request_id;
    Int arrival_time_stamp;
    Int cylinder;
    Int address;
    int process_id;
}
```
  - A file system uses contiguous allocation of disk space to files. A few blocks on the disk are reserved as spare blocks. If some disk blocks is found to be bad, the file system allocates a spare disk block to it and notes the address of the bad block and its allocated spare block in a “bad blocks table”. This table is consulted while accessing the disk block. Simulate the same.

### Additional Exercises:

- Consider the following snapshot of the system. Write C program to implement deadlock detection algorithm.
  - Is the system in a safe state?
  - Suppose that process P2 make one additional request for instance of type C, can the system still be in a safe state?

	<u>Allocation</u>	<u>Request</u>	<u>Available</u>
	<u>A B C</u>	<u>A B C</u>	<u>A B C</u>
$P_0$	0 1 0	0 0 0	0 0 0
$P_1$	2 0 0	2 0 2	
$P_2$	3 0 3	0 0 0	
$P_3$	2 1 1	1 0 0	
$P_4$	0 0 2	0 0 2	

- Display the list of devices connected to your system including the physical names and its instance number.

## WORKING WITH DIRECTORY STRUCTURES

### Objectives:

In this lab, the student will be able to:

- Understand how programs can manipulate directories
- Work with system calls to create, scan and work with directories

All subdirectory and file names within a directory must be unique. However, names within different directories can be the same. For example, the directory /usr contains the subdirectory /usr/lib. There is no conflict between /usr/lib and /lib because the path names are different.

Pathnames for files work exactly like path names for directories. The pathname of a file describes that file's place within the file system hierarchy. For example, if the /home/user2 directory contains a file called report5, the pathname for this file is /home/user2/report5. This shows that the file report5 is within the directory user2, which is within the directory home, which is within the root (/) directory.

Directories can contain only subdirectories, only files, or both.

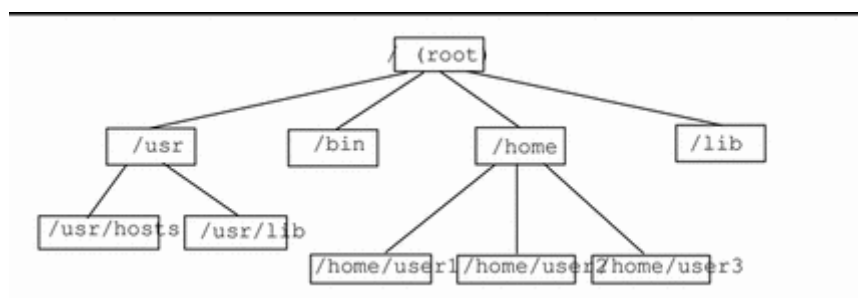


Figure 12.1 root directory

## Print Working Directory (pwd)

The pwd command will give the present working directories

```
$ pwd
/home/user1
```

## Your Home Directory

### Home

Every user has a **home** directory. When you first open the Command Tool or Shell Tool window in the OpenWindows environment, your initial location (working directory) is your home directory.

A program can determine its current working directory by calling the getcwd function.

```
#include <unistd.h>
```

```
char *getcwd(char *buf, size_t size);
```

The getcwd function writes the name of the current directory into the given buffer, buf. It returns NULL



if the directory name would exceed the size of the buffer (an ERANGE error), given as the parameter size .

\*It returns buf on success.

\*getcwd may also return NULL if the directory is removed ( EINVAL ) or permissions changed ( EACCESS ) while the program is running.

### **mkdir and rmdir**

You can create and remove directories using the mkdir and rmdir system calls.

```
#include <sys/types.h>
```

```
#include <sys/stat.h>
```

```
int mkdir(const char *path, mode_t mode);
```

The mkdir system call is used for creating directories and is the equivalent of the mkdir program. **mkdir** makes a new directory with a path as its name. The directory permissions are passed in the parameter mode and are given as in the O\_CREAT option of the open system call and, again, subject to umask .

```
#include <unistd.h>
```

```
int rmdir(const char *path);
```

The rmdir system call removes directories, but only if they are empty. The rmdir program uses this system call to do its job.

### **Scanning Directories**

A common problem on Linux systems is scanning directories, that is, determining the files that reside in a particular directory. In shell programs, it's easy — just let the shell expand a wildcard expression. In the past, different UNIX variants have allowed programmatic access to the low-level file system structure. You can still open a directory as a regular file and directly read the directory entries, but different file system structures and implementations have made this approach nonportable. A standard suite of library functions has now been developed that makes directory scanning much simpler. The directory functions are declared in a header file dirent.h. They use a structure, DIR , as a basis for directory manipulation. A pointer to this structure, called a directory stream (a DIR \* ), acts in much the same way as a file stream ( FILE \* ) does for regular file manipulation. Directory entries themselves are returned in dirent structures, also declared in dirent.h, because one should never alter the fields in the DIR structure directly.

We'll review these functions:

opendir

closedir

readdir

telldir

seekdir

closedir

### **opendir**

The opendir function opens a directory and establishes a directory stream. If successful, it returns a pointer to a DIR structure to be used for reading directory entries.

```
#include <sys/types.h>
```

```
#include <dirent.h>
```

```
DIR *opendir(const char *name);
```

opendir returns a null pointer on failure. Note that a directory stream uses a low-level file descriptor to access the directory itself, so opendir could fail with too many open files.

### **readdir**

The readdir function returns a pointer to a structure detailing the next directory entry in the directory stream dirp. Successive calls to readdir return further directory entries. On error, and at the end of the directory, readdir returns NULL. POSIX-compliant systems leave errno unchanged when returning NULL at end of directory and set it when an error occurs.

```
#include <sys/types.h>
```

```
#include <dirent.h>
```

```
struct dirent *readdir(DIR *dirp);
```

Note that readdir scanning isn't guaranteed to list all the files (and subdirectories) in a directory if other processes are creating and deleting files in the directory at the same time.

The dirent structure containing directory entry details includes the following entries:

❑ ino\_t d\_ino : The inode of the file

❑ char

d\_name[] : The name of the file

### **telldir**

The telldir function returns a value that records the current position in a directory stream. You can use

this in subsequent calls to seekdir to reset a directory scan to the current position.

```
#include <sys/types.h>
```

```
#include <dirent.h>
```

```
long int telldir(DIR *dirp);
```

### **seekdir**

The seekdir function sets the directory entry pointer in the directory stream given by dirp. The value of loc, used to set the position, should have been obtained from a prior call to telldir.

```
#include <sys/types.h>
```

```
#include <dirent.h>
```

```
void seekdir(DIR *dirp, long int loc);
```

### **closedir**

The closedir function closes a directory stream and frees up the resources associated with it. It returns 0 on success and -1 if there is an error.

```
#include <sys/types.h>
```

```
#include <dirent.h>
```

```
int closedir(DIR *dirp);
```

## **Sample Program**

A Directory-Scanning Program

1. Start with the appropriate headers and then a function, printdir, which prints out the current directory. It will recurse for subdirectories using the depth parameter for indentation.

```

#include<unistd.h>
#include<stdio.h>
#include<dirent.h>
#include<string.h>
#include<sys/stat.h>
#include<stdlib.h>

void printdir(char *dir, int depth)
{
    DIR *dp;
    struct dirent *entry;
    struct stat statbuf;
    if((dp = opendir(dir)) == NULL) {
        fprintf(stderr,"cannot open directory: %s\n", dir);
        return;
    }
    chdir(dir);
    while((entry = readdir(dp)) != NULL) {
        lstat(entry->d_name,&statbuf);
        if(S_ISDIR(statbuf.st_mode)) {
            /* Found a directory, but ignore . and .. */
            if(strcmp(".",entry->d_name) == 0 ||
               strcmp("..",entry->d_name) == 0)
                continue;
            printf("%s%s\n",depth,"",entry->d_name);
            /* Recurse at a new indent level */
            printdir(entry->d_name,depth+4);
        }
        else printf("%s%s\n",depth,"",entry->d_name);
    }
    chdir("..");
    closedir(dp);
}

```

### Lab Exercises:

1. Write a C program to emulate the `ls -l` UNIX command that prints all files in a current directory and lists access privileges, etc. DO NOT simply exec `ls -l` from the program.
2. Write a program that will list all files in a current directory and all files in subsequent subdirectories.
3. How do you list all installed programs in Linux?
4. How do you find out what RPM packages are installed on Linux?

### Additional Exercises:

1. Write a program that will only list subdirectories in alphabetical order.
2. Write a program that allows the user to remove any or all of the files in a current working directory. The name of the file should appear followed by a prompt as to whether it should be removed.

## REFERENCES:

1. Operating Systems - Internals and Design Principles", William Stallings, Prentice Hall, 5th Edition, 2004.
2. Maurice Bach, The Design of the Unix Operating System, PHI, 1986.
3. Beginning Linux Programming, 4<sup>th</sup> Edition, Neil Mathew and Richard Stores
4. Graham Glass and King Abels, Unix for Programmers and Users – *A complete guide*, PHI, 1993.
5. Sumitabha Das, Unix Concepts and Applications, McGraw Hill, 4<sup>th</sup> Edition, 2015.
6. Neil Matthew and Richard Stones, Beginning Linux Programming, 3<sup>rd</sup> Edition, Wiley, 1999.
7. A. Silberschatz, P. B. Galvin and G. Gagne, Operating System Concepts, Wiley, 8<sup>th</sup> Edition, 2014.
8. Darryl Gove, Multicore Application Programming for Windows, Linux and Oracle Solaris, Addison Wesley, 2011.
9. W. R. Stevens, UNIX Network Programming-Volume II (IPC), PHI, 1998.